



```
import pandas as pd
import numpy as np
import seaborn as sns
import math
import torch
import torch.nn as nn
import torch.optim as optim
import pytorch_lightning as pl
from pytorch_lightning.callbacks import ModelCheckpoint, EarlyStopping
from torch.utils.data import Dataset, DataLoader
from sklearn.preprocessing import MinMaxScaler
import matplotlib.pyplot as plt
%matplotlib inline
from helpers import *
import tqdm as tq
```

```
pl.seed_everything(123)
import yfinance as yf
tq.tqdm.pandas()
```

INFO:lightning_fabric.utilities.seed:Global seed set to 123

compile_stocks()

- Yahoo Finance only allows 7 days worth of data at 1 minute intervals
- This function will compile dataframes in 7 day increments for an entire range of dates

- They also only allow 1 min data for up to 30 days prior
- This can be used for other intervals that allow for longer time periods

```
def compile_stocks(symbol, end, start, day_window, interval):
    import datetime
    import yfinance as yf

    end_date = end
    start_date = (pd.to_datetime(end) - datetime.timedelta(days = day_window))

    dfs = []
    stop_me = False

    while pd.to_datetime(start_date) >= pd.to_datetime(start):
        df = yf.download(symbol,
                        start = start_date,
                        end = end_date,
                        interval = interval)

        dfs.append(df)
        end_date = start_date
        start_date = start_date - datetime.timedelta(days = day_window)

        if start_date < pd.to_datetime(start):
            start_date = pd.to_datetime(start)
        else:
            start_date = start_date

        if start_date == end_date:
            break

    master_df = pd.concat(dfs).sort_values(by="Datetime")

    return master_df
```

```
data = compile_stocks('ETH-USD',
                    '2023-01-24',
                    '2022-12-27',
                    7, '1m')
```

```
[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed
```

The Data

```
df = data[['Close', 'Volume']].copy()
df.columns = [x.lower() for x in df.columns]
head_tail_horz(df, 3, 'Raw Ethereum Data', intraday=True)
```

Raw Ethereum Data

head(3)			tail(3)		
	close	volume		close	volume
Datetime			Datetime		
2022-12-27 00:00:00+00:00	1,226.99	0	2023-01-23 23:56:00+00:00	1,629.03	323,584
2022-12-27 00:01:00+00:00	1,227.44	0	2023-01-23 23:57:00+00:00	1,628.99	716,288
2022-12-27 00:02:00+00:00	1,228.57	14,410,752	2023-01-23 23:58:00+00:00	1,628.44	0

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
DatetimeIndex: 40100 entries, 2022-12-27 00:00:00+00:00 to 2023-01-23 23:58:00+00:00
```

```
Data columns (total 2 columns):
```

```
#   Column  Non-Null Count  Dtype
---  -
0   close   40100 non-null    float64
1   volume  40100 non-null    int64
```

```
dtypes: float64(1), int64(1)
```

```
memory usage: 939.8 KB
```

```
df.shape
```

```
(40100, 2)
```

SECTIONS: [Top](#) | [Compile Stocks](#) | [The Data](#) | [Add Change Column](#) | [Feature Engineering](#) | [Splitting & Scaling](#) | [Creating Sequences](#) | [PyTorch Datasets](#) | [LSTM](#) | [Trainer](#) | [Results](#) |

add_change_column()

Shift Method

- using `shift()` to add the previous timestamp value
- then creating a difference column to show change since last
- `add_change_column()` is the function that will do all this

```
def add_change_column(df, column_changing, new_col_name):
    df['previous'] = df[column_changing].shift()
    df = df.drop(df.index[0])
    df[new_col_name] = df[column_changing] - df.previous
    df = df.drop(columns = ['previous'])
    return df
```

```
df = add_change_column(df, 'close', 'change')
```

```
head_tail_vert(df, 3, "Change column added", intraday=True)
```

Change column added: head(3)

	close	volume	change
Datetime			
2022-12-27 00:01:00+00:00	1,227.44	0	0.45
2022-12-27 00:02:00+00:00	1,228.57	14410752	1.13
2022-12-27 00:03:00+00:00	1,228.84	6625536	0.27

Change column added: tail(3)

	close	volume	change
Datetime			
2023-01-23 23:56:00+00:00	1,629.03	323584	0.14
2023-01-23 23:57:00+00:00	1,628.99	716288	-0.03
2023-01-23 23:58:00+00:00	1,628.44	0	-0.55

SECTIONS: [Top](#) | [Compile Stocks](#) | [The Data](#) | [Add Change Column](#) | [Feature Engineering](#) | [Splitting & Scaling](#) | [Creating Sequences](#) | [PyTorch Datasets](#) | [LSTM](#) | [Trainer](#) | [Results](#) |

featurize_stocks() - Feature Engineering

```
def featurize_stocks(df):
    df['weekday'] = df.index.dayofweek
    df['month_day'] = df.index.day
    df['year_week'] = df.index.isocalendar().week
    df['month'] = df.index.month

    return df
```

```
df = featurize_stocks(df)
```

```
head_tail_vert(df, 3, 'df with added features')
```

df with added features: head(3)

	close	volume	change	weekday	month_day	year_week	month
Datetime							
2022-12-27	1,227.44	0	0.45	1	27	52	12
2022-12-27	1,228.57	14410752	1.13	1	27	52	12
2022-12-27	1,228.84	6625536	0.27	1	27	52	12

df with added features: tail(3)

	close	volume	change	weekday	month_day	year_week	month
Datetime							
2023-01-23	1,629.03	323584	0.14	0	23	4	1
2023-01-23	1,628.99	716288	-0.03	0	23	4	1
2023-01-23	1,628.44	0	-0.55	0	23	4	1

```
describe_em(df, ['close', 'volume', 'change'])
```

df.close		df.volume		df.change	
	close		volume		change
count	40,099.00	count	40,099.00	count	40,099.00
mean	1,375.36	mean	2,073,010.21	mean	0.01
std	164.02	std	8,437,534.17	std	0.51
min	1,185.70	min	0.00	min	-22.17
25%	1,217.39	25%	0.00	25%	-0.12
50%	1,324.79	50%	167,424.00	50%	0.00
75%	1,548.57	75%	1,481,728.00	75%	0.13
max	1,674.18	max	618,774,528.00	max	29.07

```
df.shape
```

```
(40099, 7)
```

Splitting & Scaling the Data

```
train_size = int(len(df) * .8)
pretty(f'number of training inputs: {train_size:,}.')
test_size = int(len(df) * .2)
pretty(f'number of testing inputs: {test_size:,}.')
```

number of training inputs: 32,079.

number of testing inputs: 8,019.

```
train_df, test_df = df[:train_size], df[train_size + 1:]
```

```
pretty(f'train_df.shape: {train_df.shape} | test_df.shape: {test_df.shape}')
```

train_df.shape: (32079, 7) | test_df.shape: (8019, 7)

Scaling the data

```
scaler = MinMaxScaler(feature_range = (-1, 1))
scaler = scaler.fit(train_df)
```

```
train_df = pd.DataFrame(scaler.transform(train_df),
                        index = train_df.index,
                        columns = train_df.columns)

test_df = pd.DataFrame(scaler.transform(test_df),
                      index = test_df.index,
                      columns = test_df.columns)
```

Viewing the scaled data

```
multi([(train_df.head(5), 'train_df.head(5)'),
      (train_df.tail(5), 'train_df.tail(5)'),
      (test_df.head(5), 'test_df.head(5)'),
      (test_df.tail(5), 'test_df.tail(5)')], intraday=True)
```

train_df.head(5)							
	close	volume	change	weekday	month_day	year_week	month
Datetime							
2022-12-27 00:01:00+00:00	-0.80	-1.00	-0.50	-0.67	0.73	1.00	1.00
2022-12-27 00:02:00+00:00	-0.79	-0.94	-0.46	-0.67	0.73	1.00	1.00
2022-12-27 00:03:00+00:00	-0.79	-0.97	-0.50	-0.67	0.73	1.00	1.00
2022-12-27 00:04:00+00:00	-0.78	-0.88	-0.44	-0.67	0.73	1.00	1.00
2022-12-27 00:05:00+00:00	-0.78	-0.99	-0.52	-0.67	0.73	1.00	1.00

train_df.tail(5)							
	close	volume	change	weekday	month_day	year_week	month
Datetime							
2023-01-18 09:16:00+00:00	0.92	-0.97	-0.50	-0.33	0.13	-0.92	-1.00
2023-01-18 09:17:00+00:00	0.92	-1.00	-0.52	-0.33	0.13	-0.92	-1.00
2023-01-18 09:18:00+00:00	0.92	-1.00	-0.53	-0.33	0.13	-0.92	-1.00
2023-01-18 09:19:00+00:00	0.92	-0.99	-0.50	-0.33	0.13	-0.92	-1.00
2023-01-18 09:20:00+00:00	0.93	-0.99	-0.49	-0.33	0.13	-0.92	-1.00

test_df.head(5)							
	close	volume	change	weekday	month_day	year_week	month
Datetime							
2023-01-18 09:22:00+00:00	0.93	-0.98	-0.47	-0.33	0.13	-0.92	-1.00
2023-01-18 09:23:00+00:00	0.93	-1.00	-0.50	-0.33	0.13	-0.92	-1.00
2023-01-18 09:24:00+00:00	0.93	-1.00	-0.51	-0.33	0.13	-0.92	-1.00
2023-01-18 09:25:00+00:00	0.93	-0.99	-0.51	-0.33	0.13	-0.92	-1.00
2023-01-18 09:26:00+00:00	0.93	-1.00	-0.52	-0.33	0.13	-0.92	-1.00

test_df.tail(5)							
	close	volume	change	weekday	month_day	year_week	month
Datetime							
2023-01-23 23:54:00+00:00	1.17	-1.00	-0.52	-1.00	0.47	-0.88	-1.00
2023-01-23 23:55:00+00:00	1.17	-0.99	-0.51	-1.00	0.47	-0.88	-1.00
2023-01-23 23:56:00+00:00	1.17	-1.00	-0.51	-1.00	0.47	-0.88	-1.00
2023-01-23 23:57:00+00:00	1.17	-1.00	-0.52	-1.00	0.47	-0.88	-1.00
2023-01-23 23:58:00+00:00	1.17	-1.00	-0.55	-1.00	0.47	-0.88	-1.00

SECTIONS: [Top](#) | [Compile Stocks](#) | [The Data](#) | [Add Change Column](#) | [Feature Engineering](#) | [Splitting & Scaling](#) | [Creating Sequences](#) | [PyTorch Datasets](#) | [LSTM](#) | [Trainer](#) | [Results](#) |

Creating sequences for training the model

compile_sequences()

- will take input data as dataframe, a target column and the length we want the sequences to be
 - This will iterate over all the data from the beginning up to the last possible full sequence length in the data.
 - It is also possible to write this code in a more elegant way, such as with sliding windows

compile_sequences()

```
def compile_sequences(input_data, target_column, sequence_length):

    sequences = []
    data_size = len(input_data)

    for item in tqdm(range(data_size - sequence_length)):

        # end not included, therefor it will be the label
        sequence = input_data[item : item + sequence_length]
        label_position = item + sequence_length

        # defining the label with the value in the label index position
        label = input_data.iloc[label_position][target_column]
        sequences.append((sequence, label))

    return sequences
```

How creating sequences works

```
sample_data = pd.DataFrame(dict(
    feature1 = [1, 2, 3, 4, 5, 6, 7, 8],
    label = [6, 7, 8, 9, 10, 11, 12, 13]))
```

```
sample_data.head(5)
```

	feature1	label
0	1	6
1	2	7
2	3	8
3	4	9
4	5	10

```
sample_sequences = compile_sequences(sample_data, 'label', 3)
```

```
100%|██████████| 5/5 [00:00<00:00, 1095.58it/s]
```

see_sequence_samples()

```
def see_sequence_samples(data, feature_list, num_samples,
                        sequence_portion=5):

    total_sequences = len(data)

    div_print(f'There are {total_sequences:,} total sequences in this data.',
              fontsize = 4)
    counter = 1

    for record in range(0, num_samples):
```



```

see(data[record][0][feature_list].head(sequence_portion),
    f'sequence no.{record + 1} | target / label -> {data[record][1]}',
    fontsize = 3)
counter += 1

```

```

see_sequence_samples(sample_sequences, ['feature1'], num_samples = 3)

```

There are 5 total sequences in this data.

sequence no.1 | target / label -> 9

feature1	
0	1
1	2
2	3

sequence no.2 | target / label -> 10

feature1	
1	2
2	3
3	4

sequence no.3 | target / label -> 11

feature1	
2	3
3	4
4	5

Creating sequences from the data

```

sequence_length = 60

train_sequences = compile_sequences(train_df,
                                   'close',
                                   sequence_length)

test_sequences = compile_sequences(test_df,
                                   'close',
                                   sequence_length)

```

100%|██████████| 32019/32019 [00:12<00:00, 2618.53it/s]

100%|██████████| 7959/7959 [00:03<00:00, 2364.64it/s]

```
see_sequence_samples(train_sequences,
                     ['close', 'volume', 'change'],
                     num_samples = 2,
                     sequence_portion = 3)
```

There are 32,019 total squences in this data.

sequence no.1 | target / label -> -0.8090204723204248

	close	volume	change
Datetime			
2022-12-27	-0.80	-1.00	-0.50
2022-12-27	-0.79	-0.94	-0.46
2022-12-27	-0.79	-0.97	-0.50

sequence no.2 | target / label -> -0.8085050919224583

	close	volume	change
Datetime			
2022-12-27	-0.79	-0.94	-0.46
2022-12-27	-0.79	-0.97	-0.50
2022-12-27	-0.78	-0.88	-0.44

```
see_sequence_samples(test_sequences,
                     ['close', 'volume', 'change'],
                     num_samples = 2,
                     sequence_portion = 3)
```

There are 7,959 total squences in this data.

sequence no.1 | target / label -> 0.9247891394642078

	close	volume	change
Datetime			
2023-01-18	0.93	-0.98	-0.47
2023-01-18	0.93	-1.00	-0.50
2023-01-18	0.93	-1.00	-0.51

sequence no.2 | target / label -> 0.9241763031441197

	close	volume	change
Datetime			
2023-01-18	0.93	-1.00	-0.50
2023-01-18	0.93	-1.00	-0.51
2023-01-18	0.93	-0.99	-0.51

Sequences to PyTorch Datasets

- `__init__()` - takes the sequences it will work with
- `__len__()` - returns the number of sequences in any given example
- `__getitem__()` - takes the index of the item we are interested in
 - within each sequence item is a tuple with the input data as the first tuple value and the target data or label as the second
 - sequences are converted from pandas to numpy to Tensor
 - labels are converted into floats, since we are doing regression and wanting to predict floating point numbers

```
class TickerDataset(Dataset):
    def __init__(self, sequences):
        self.sequences = sequences

    def __len__(self):
        return len(self.sequences)

    def __getitem__(self, idx):
        sequence, label = self.sequences[idx]
        return dict(sequence = torch.Tensor(sequence.to_numpy()),
                    label = torch.tensor(label).float())
```

Wrapping data with PyTorch Lightning

- `setup()` converts sequences into the dataset class defined above
- `train_dataloader(), validation_dataloader(), test_dataloader()`
 - create the three different kinds of dataloaders for the data
 - `train_dataloader()` -> `shuffle = False` - do not want to shuffle the data because it is timeseries data, so the order is important
 - `num_workers = 2` - speeds up training depending on GPU vs CPU
 - `validation_dataloader()` and `test_dataloader()` - `batch_size = 1` - for testing and making predictions, we would usually do one record at a time

```
class TickerDataModule(pl.LightningDataModule):

    def __init__(self,
                 train_sequences,
                 test_sequences,
```

```

        batch_size = 8):
    super().__init__()
    self.train_sequences = train_sequences
    self.test_sequences = test_sequences
    self.batch_size = batch_size

    def setup(self, stage = None):
        self.train_dataset = TickerDataset(self.train_sequences)
        self.test_dataset = TickerDataset(self.test_sequences)

    def train_dataloader(self):
        return DataLoader(self.train_dataset,
                          batch_size = self.batch_size,
                          shuffle = False,
                          num_workers = 2)

    def validation_dataloader(self):
        return DataLoader(self.test_dataset,
                          batch_size = 1,
                          shuffle = False,
                          num_workers = 1)

    def test_dataloader(self):
        return DataLoader(self.test_dataset,
                          batch_size = 1,
                          shuffle = False,
                          num_workers = 1)

```

Defining the data module

- calling the module with the sequences, with `setup()`, which will create the PyTorch Datasets

```

batch_size = 64

data_module = TickerDataModule(train_sequences,
                               test_sequences,
                               batch_size = batch_size)

data_module.setup(stage = None)

```

Investigating a single item from TickerDataset

```

train_dataset = TickerDataset(train_sequences)

```

```

for item in train_dataset:
    print(item['sequence'].shape)
    print(item['label'].shape)
    print(item['label'])
    break

```

```
torch.Size([60, 7])
torch.Size([1])
tensor(-0.8090)
```

SECTIONS: [Top](#) | [Compile Stocks](#) | [The Data](#) | [Add Change Column](#) | [Feature Engineering](#) | [Splitting & Scaling](#) | [Creating Sequences](#) | [PyTorch Datasets](#) | [LSTM](#) | [Trainer](#) | [Results](#) |

Creating the LSTM model

- LSTM = Long short term memory neural network
 - Allows sequences with single outputs or multiple datapoints
 - Transformers can also be used for time series (as used in text and images)
 - `batch_size = True` - inputting batch size as first parameter, which is useful for using the dataloader that passes batches of sequences with 120 data points in each sequence and 9 features for each data point
 - `num_layers` - number of LSTMs stacked on one another
 - regressor is the final linear layer that outputs our single prediction
- `forward()`
 - `flatten_parameters()` - used for quicker distributed training with GPUs, which PyTorch lightning allows for
 - takes the features of the output of final layer of the LSTM, the hidden state and the cell state of the LSTM

```
class PricePredictionModel(nn.Module):

    def __init__(self,
                  num_features,
                  num_hidden = 128,
                  num_layers = 2):

        super().__init__()
        self.num_hidden = num_hidden

        self.lstm = nn.LSTM(
            input_size = num_features,
            hidden_size = num_hidden,
            batch_first = True,
            num_layers = num_layers,
            dropout = 0.2)

        self.regressor = nn.Linear(num_hidden, 1)

    def forward(self, inputs):
        self.lstm.flatten_parameters()

        # getting states from the LSTM, retrieving the last layer of hidden
        # passing the last layer to the regressor
```

```
_, (hidden, _) = self.lstm(inputs)
out = hidden[-1]
return self.regressor(out)
```

Creating the lightning module for the model

- `labels.unsqueeze()` - aligns the dimensions of the output with the output of predictions from the model

```
class TickerPricePredictor(pl.LightningModule):

    def __init__(self, num_features):
        super().__init__()

        self.model = PricePredictionModel(num_features)
        self.loss_function = nn.MSELoss()

    def forward(self, inputs, labels = None):
        output = self.model(inputs)
        loss = 0
        if labels is not None:
            loss = self.loss_function(output, labels.unsqueeze(dim=1))
        return loss, output

    def training_step(self, batch, batch_idx):
        sequences = batch['sequence']
        labels = batch['label']
        loss, outputs = self(sequences, labels)
        self.log('training loss', loss, prog_bar = True, logger = True)
        return loss

    def validation_step(self, batch, batch_idx):
        sequences = batch['sequence']
        labels = batch['label']
        loss, outputs = self(sequences, labels)
        self.log('validation loss', loss, prog_bar = True, logger = True)
        return loss

    def test_step(self, batch, batch_idx):
        sequences = batch['sequence']
        loss, outputs = self(sequences, labels)
        self.log('test loss', loss, prog_bar = True, logger = True)
        return loss

    def configure_optimizers(self):
        return optim.AdamW(self.parameters(), lr = 0.0001)
```

```
model = TickerPricePredictor(num_features = train_df.shape[1])
```

Inspecting the values within a single item in the training dataloader

NOTICE: If the following does not work, FYI it works in Google Colab. There is a PyTorch Lightning issue, even though the version is identical.

```
for item in data_module.train_dataloader():
    sequence_shape = item['sequence'].shape
    label_shape = item['label'].shape
    break

pretty(sequence_shape, "batches of sequences: 64 -> each sequence length: 60 -> num
pretty(label_shape, "one label for each batch of sequences: 64")
pretty(len(data_module.train_dataloader()), "total number of batches: 501 (len(dataload
```

batches of sequences: 64 -> each sequence length: 60 -> num_features for each sequence: 7

torch.Size([64, 60, 7])

one label for each batch of sequences: 64

torch.Size([64])

total number of batches: 501 (len(dataloader))

501

Checkpoint Callback

- `dirpath = 'checkpoints'` - the folder where the checkpoints will be saved
- `filename = 'best_checkpoint'` - name that the best model saved during training will be saved as
- `save_top_k = 1` - how many top models to save
- `verbose = True` - how much we want to know about what is going on during training
- `monitor = 'training loss'` - the metrics we want to monitor while training
- `mode = 'min'` - what aspect of the metric to monitor

```
checkpoints = ModelCheckpoint(dirpath = 'checkpoints',
                              filename = 'best_checkpoint',
                              save_top_k = 1,
                              verbose = True,
                              monitor = 'training loss',
                              mode = 'min')
```

- `logger = TensorBoardLogger('lightning_logs', name = 'bitcoin_price')` - establishing the logger through TensorBoard and naming the project
- `early_stopping_callback = EarlyStopping(monitor = 'val loss')` - a second callback that watches the validation loss and if there is no improvement for 2 epochs (`patience = 2`), the training will stop

```
early_stopping = EarlyStopping(monitor = 'training loss', patience = 3)
```

Training the Model

- setting up the trainer to train the model

```
num_epochs = 10

trainer = pl.Trainer(callbacks = [checkpoints, early_stopping],
                    max_epochs = num_epochs,
                    auto_lr_find=True,
                    gpus = 1)
```

```
/usr/local/lib/python3.8/dist-
packages/pytorch_lightning/trainer/connectors/accelerator_connector.py:467:
LightningDeprecationWarning: Setting `Trainer(gpus=1)` is deprecated in v1.7 and will
be removed in v2.0. Please use `Trainer(accelerator='gpu', devices=1)` instead.
    rank_zero_deprecation(
INFO:pytorch_lightning.utilities.rank_zero:GPU available: True (cuda), used: True
INFO:pytorch_lightning.utilities.rank_zero:TPU available: False, using: 0 TPU cores
INFO:pytorch_lightning.utilities.rank_zero:IPU available: False, using: 0 IPU
INFO:pytorch_lightning.utilities.rank_zero:HPU available: False, using: 0 HPUs
```

```
trainer.fit(model, data_module)
```

```
/usr/local/lib/python3.8/dist-
packages/pytorch_lightning/trainer/configuration_validator.py:108: PossibleUserWarning:
You defined a `validation_step` but have no `val_dataloader`. Skipping val loop.
    rank_zero_warn(
/usr/local/lib/python3.8/dist-
packages/pytorch_lightning/callbacks/model_checkpoint.py:612: UserWarning: Checkpoint
directory /content/checkpoints exists and is not empty.
    rank_zero_warn(f"Checkpoint directory {dirpath} exists and is not empty.")
INFO:pytorch_lightning.accelerators.cuda:LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]
INFO:pytorch_lightning.callbacks.model_summary:
  | Name                | Type                | Params
-----
0 | model                | PricePredictionModel | 202 K
1 | loss_function        | MSELoss             | 0
-----
202 K    Trainable params
0        Non-trainable params
```


202 K Total params
0.809 Total estimated model params size (MB)

Training: 0it [00:00, ?it/s]

INFO:pytorch_lightning.utilities.rank_zero:Epoch 0, global step 501: 'training loss' reached 0.00245 (best 0.00245), saving model to '/content/checkpoints/best_checkpoint-v5.ckpt' as top 1

INFO:pytorch_lightning.utilities.rank_zero:Epoch 1, global step 1002: 'training loss' reached 0.00207 (best 0.00207), saving model to '/content/checkpoints/best_checkpoint-v5.ckpt' as top 1

INFO:pytorch_lightning.utilities.rank_zero:Epoch 2, global step 1503: 'training loss' reached 0.00100 (best 0.00100), saving model to '/content/checkpoints/best_checkpoint-v5.ckpt' as top 1

INFO:pytorch_lightning.utilities.rank_zero:Epoch 3, global step 2004: 'training loss' was not in top 1

INFO:pytorch_lightning.utilities.rank_zero:Epoch 4, global step 2505: 'training loss' was not in top 1

INFO:pytorch_lightning.utilities.rank_zero:Epoch 5, global step 3006: 'training loss' was not in top 1

SECTIONS: [Top](#) | [Compile Stocks](#) | [The Data](#) | [Add Change Column](#) | [Feature Engineering](#) | [Splitting & Scaling](#) | [Creating Sequences](#) | [PyTorch Datasets](#) | [LSTM](#) | [Trainer](#) | [Results](#) |

Loading a Pre-Trained Model & Getting Predictions

```
trained_model = TickerPricePredictor.load_from_checkpoint('/content/checkpoints/best_checkpoint-v5.ckpt', num_features = train_df.shape[1])  
  
trained_model.freeze()
```

```
test_dataset = TickerDataset(test_sequences)
```

get_predictions()

```
def get_predictions(model,  
                    test_dataset,  
                    original_df = df,  
                    train_size = len(train_df),  
                    sequence_length = sequence_length,  
                    scaler = scaler,  
                    plot = True):  
  
    predictions = []  
    labels = []  
    for item in tq.tqdm(test_dataset):  
        sequence = item['sequence']
```

```

label = item['label']

_, output = model(sequence.unsqueeze(dim=0))
predictions.append(output.item())
labels.append(label.item())

test_data = df[train_size + 1 : ]
test_sequences_data = test_data.iloc[sequence_length:]
dates = test_sequences_data.index.tolist()

results = pd.concat([pd.Series(dates), pd.Series(predictions),
                        pd.Series(labels)], axis = 1)
results.columns = ['dates', 'predictions', 'actuals']
results = results.set_index('dates')

def reverse_scale(scaler, data):
    reverse_scaler = MinMaxScaler()
    reverse_scaler.min_ = scaler.min_[-1]
    reverse_scaler.scale_ = scaler.scale_[-1]
    data_array = np.array(data)[: , np.newaxis]
    results = reverse_scaler.inverse_transform(data_array).flatten()
    return results

results.predictions = reverse_scale(scaler, results.predictions)
results.actuals = reverse_scale(scaler, results.actuals)

results['accuracy'] = (1 - abs(results.actuals - results.predictions) /
                        results.actuals) * 100

overall_test_accuracy = results.accuracy.mean()

sp(); sp();

pretty(f'{overall_test_accuracy:.3f}%',
       'The overall testing accuracy of the model is:'); sp()

if plot == True:
    fig = plt.subplots(facecolor='#222222', figsize=(13, 7))
    ax = plt.axes();
    plt.style.use("ggplot");
    ax.set_facecolor('#333333')
    ax.grid(color=fontcolor, linestyle=':', linewidth=0.75, alpha=0.75)
    plt.tick_params(labelrotation=40);
    plt.title('Predictions vs. Actual Labels', fontsize=23, pad=20, color=fontcolor)
    plt.ylabel('Stock Prices', fontsize=18, color=fontcolor);
    plt.xlabel('Test Sequences', fontsize=18, color=fontcolor);
    plt.xticks(fontsize=10, color='white')
    plt.yticks(fontsize=10, color='white')
    results.predictions.plot(ax = ax, color = 'cyan');
    results.actuals.plot(ax = ax, color = 'deeppink');

```

```
plt.legend(facecolor = 'DarkGray', loc = 2, fontsize=15);
```

```
return results
```

```
results = get_predictions(model, test_dataset)
```

100%|██████████| 7959/7959 [01:29<00:00, 88.59it/s]

The overall testing accuracy of the model is:

96.367%

