

---

Sections: ● [Top](#) ● [The Data](#) ● [Feature Engineering](#) ● [Investigating Correlation](#) ● [Lag Features](#) ● [Splitting](#) ● [The Model](#) ● [Results with Traditional Split](#) ● [Using Cross-Validation](#) ● [Making Future Predictions](#)

---

# Time Series Prediction with **XGBoost** ~ EVAN MARIE CARIZ

Evan Marie online: [EvanMarie@Proton.me](mailto:EvanMarie@Proton.me) | [Linked In](#) | [GitHub](#) | [Hugging Face](#) | [Mastadon](#)  
| [Jovian.ai](#) | [TikTok](#) | [CodeWars](#) | Discord ⇒ ☆☆ EvanMarie ☆☆ #6114

```
from helpers import *
import_all()
from xgboost import XGBRegressor
%matplotlib inline
import seaborn as sns
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import TimeSeriesSplit
```

---

Sections: ● [Top](#) ● [The Data](#) ● [Feature Engineering](#) ● [Investigating Correlation](#) ● [Lag Features](#) ● [Splitting](#) ● [The Model](#) ● [Results with Traditional Split](#) ● [Using Cross-Validation](#) ● [Making Future Predictions](#)

---

## The Data

- This data is an excerpt from a [Kaggle maintained and regularly update dataset collection](#)
- The dataset reflects the energy consumption as reported by the National Grid ESO, Great Britain's electricity system operator
- Consumption is recorded twice an hour
- The data covers January 1, 2009 to December 31, 2022

## Importing Data

```
data = pd.read_csv('uk_power_consumption.csv', parse_dates = ['settlement_date'])
data = data[['settlement_date', 'tsd', 'is_holiday']]
data.columns = ['datetime', 'consumption', 'holiday']
data = data.set_index('datetime', drop=True)
```

```
head_tail_horz(data, 5, "UK Power Consumption Data", intraday = True)
```

### UK Power Consumption Data

head(5)			tail(5)		
consumption		holiday	consumption		holiday
datetime			datetime		
2009-01-01 00:00:00	38,704	1	2022-12-31 21:30:00	25,634	0
2009-01-01 00:30:00	38,964	1	2022-12-31 22:00:00	24,788	0
2009-01-01 01:00:00	38,651	1	2022-12-31 22:30:00	24,365	0
2009-01-01 01:30:00	37,775	1	2022-12-31 23:00:00	24,766	0
2009-01-01 02:00:00	37,298	1	2022-12-31 23:30:00	24,843	0

```
def timeseries_overview(df, main_col):
    index_col = ['number of records', 'number of columns',
                 'missing values', 'columns', 'start date',
                 'end date', 'main column min', 'main column max']

    num_records = df.shape[0]
    num_cols = df.shape[1]
    missing_values = df.isna().sum().sum()
    columns = ', '.join(list(df.columns))
    start_date = min(df.index).strftime('%m/%d/%Y')
    end_date = max(df.index).strftime('%m/%d/%Y')
    main_min = min(df[main_col])
    main_max = max(df[main_col])

    values = [num_records, num_cols, missing_values,
              columns, start_date, end_date, main_min, main_max]

    overview = pd.concat([pd.Series(index_col), pd.Series(values)], axis = 1)
    overview.columns = ['aspects', 'information']
    overview.set_index('aspects')

    styling = {'information': [{'selector': '',
                                'props': [('font-size', '15px'),
                                           ('padding-right', '15px'),
                                           ('padding-left', '35px')]}],
               'aspects': [{'selector': '',
                             'props': [('font-weight', 'bold'),
                                        ('font-size', '15px'),
                                        ('padding-right', '15px'),
                                        ('padding-left', '15px')]}]}

    pretty('Initial DataFrame Overview', fontsize=4)

    return overview.style\
```

```
.hide(axis='index')\
.set_table_styles(styling)\
.format(precision=0, thousands=",")
```

```
timeseries_overview(data, 'consumption')
```

## Initial DataFrame Overview


aspects		information
number of records		244,676
number of columns		2
missing values		0
columns	consumption, holiday	
start date		01/01/2009
end date		12/31/2022
main column min		16,629
main column max		60,147

## Initial Plotting

```
colors = sns.color_palette('husl', 11)
colors
```




```
fig = plt.figure(facecolor = '#333333', figsize = (13,7));
ax = plt.axes();
ax.set_facecolor('#444444')
data['consumption'].plot(color = colors[9],
                        ax = ax);
ax.grid()
plt.xlabel('years', color = 'white', fontsize = 15)
plt.ylabel('energy usage', color = 'white', fontsize = 15)
plt.xticks(color='white'); plt.yticks(color='white');
plt.title('Energy Consumption in the UK: 2009-2022', fontsize = 20,
        pad = 20, color = 'white');
```




## Looking at a Single Week of Data

```
fig, ax = plt.subplots(figsize = (13, 7), facecolor = '#333333')
ax.set_facecolor('#444444')
data['consumption'].iloc[(data.index > '01-23-2013') \
                        & (data.index < '01-30-2013')].plot(ax = ax, linewidth = 3, color = 'cyan')
ax.grid()
plt.xlabel('dates', color = 'white', fontsize = 15)
plt.ylabel('energy usage', color = 'white', fontsize = 15)
plt.tick_params(labelcolor = 'white', which = 'both')
plt.title('Energy Usage: January 23-30, 2013', fontsize = 20,
          pad = 20, color = 'white');
plt.legend().remove()
```



## Distribution of Energy Use Values

```
fig, ax = plt.subplots(figsize = (13, 7), facecolor = '#333333')
ax.set_facecolor('#444444')
data['consumption'].plot(kind = 'hist', bins = 100,
                        color = 'deeppink', edgecolor = 'white');
ax.grid()
plt.xlabel('dates', color = 'white', fontsize = 15)
plt.ylabel('energy usage', color = 'white', fontsize = 15)
plt.tick_params(labelcolor = 'white', which = 'both')
plt.title('Distribution of Energy Use Values',
        fontsize = 20, pad = 20, color = 'white');
plt.legend().remove()
```



Sections: ● [Top](#) ● [The Data](#) ● [Feature Engineering](#) ● [Investigating Correlation](#) ● [Lag Features](#) ● [Splitting](#) ● [The Model](#) ● [Results with Traditional Split](#) ● [Using Cross-Validation](#) ● [Making Future Predictions](#)

## Feature Engineering

### Time Series DateTime Index Feature Creation

```
def featurize_datetime_index(df, daytime = True):
    df = df.copy()

    df['hour'] = df.index.hour
    df['weekday'] = df.index.dayofweek
    df['weekday_name'] = df.index.strftime('%A')
    df['month'] = df.index.month
    df['month_name'] = df.index.strftime('%B')
    df['quarter'] = df.index.quarter
    df['year'] = df.index.year
    df['week_of_year'] = df.index.isocalendar().week
    df['day_of_year'] = df.index.dayofyear

    if daytime:
        # Add column with category for time of day:
        # midnight, early_morning, late_morning, afternoon, evening, night
        def time_of_day(hour):
            if hour >= 0 and hour < 6:
                return 'midnight'
            elif hour >= 6 and hour < 9:
                return 'early_morning'
```

```

elif hour >= 9 and hour < 12:
    return 'late_morning'
elif hour >= 12 and hour < 15:
    return 'afternoon'
elif hour >= 15 and hour < 18:
    return 'evening'
else:
    return 'night'

df['time_of_day'] = (df['hour'].apply(time_of_day)).astype('category')

df['weekday_name'] = df['weekday_name'].astype('category')
df['month_name'] = df['month_name'].astype('category')

return df

```

```
df = featurize_datetime_index(data.copy())
```

```
head_tail_horz(df.sample(10), 5, 'DF with Added Datetime Features (Random Samples)', in
```

#### DF with Added Datetime Features (Random Samples)

head(5)

	consumption	holiday	hour	weekday	weekday_name	month	month_name	quarter	year	week_of_year	date
datetime											
2019-03-23 18:00:00	37,097	0	18	5	Saturday	3	March	1	2,019	12	
2021-10-07 21:00:00	30,872	0	21	3	Thursday	10	October	4	2,021	40	
2022-06-20 08:30:00	30,410	0	8	0	Monday	6	June	2	2,022	25	
2010-11-15 23:00:00	36,683	0	23	0	Monday	11	November	4	2,010	46	
2016-11-06 21:00:00	35,070	0	21	6	Sunday	11	November	4	2,016	44	

tail(5)

	consumption	holiday	hour	weekday	weekday_name	month	month_name	quarter	year	week_of_year	date
datetime											
2014-05-08 14:30:00	39,753	0	14	3	Thursday	5	May	2	2,014	19	
2011-04-05 12:30:00	43,494	0	12	1	Tuesday	4	April	2	2,011	14	
2011-12-01 06:00:00	33,868	0	6	3	Thursday	12	December	4	2,011	48	
2012-04-03 07:30:00	38,705	0	7	1	Tuesday	4	April	2	2,012	14	
2019-05-04 17:00:00	26,747	0	17	5	Saturday	5	May	2	2,019	18	

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
DatetimeIndex: 244676 entries, 2009-01-01 00:00:00 to 2022-12-31 23:30:00
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	consumption	244676 non-null	int64
1	holiday	244676 non-null	int64
2	hour	244676 non-null	int64
3	weekday	244676 non-null	int64
4	weekday_name	244676 non-null	category
5	month	244676 non-null	int64
6	month_name	244676 non-null	category
7	quarter	244676 non-null	int64
8	year	244676 non-null	int64
9	week_of_year	244676 non-null	UInt32
10	day_of_year	244676 non-null	int64
11	time_of_day	244676 non-null	category

```
dtypes: UInt32(1), category(3), int64(8)
```

```
memory usage: 18.7 MB
```




## Investigating Correlation

### Visualize Correlation Between Features and Target


```
def boxplot_correlation(df, feature_x, feature_y, order = None, palette = None):  
  
    fig, ax = plt.subplots(figsize = (13, 7), facecolor = '#333333')  
    ax.set_facecolor('LightGray')  
  
    sns.boxplot(data = df,  
                x = feature_x,  
                y = feature_y,  
                order = order,  
                palette = palette)  
  
    x_name = str(df[feature_x].name)  
    y_name = str(df[feature_y].name)  
  
    ax.grid()  
    plt.xlabel(x_name, color = 'white', fontsize = 15)  
    plt.ylabel(y_name, color = 'white', fontsize = 15)  
    plt.xticks(color='white'); plt.yticks(color='white');  
    plt.title(f'Feature Correlation: {x_name.capitalize()} - {y_name.capitalize()}',  
            fontsize = 20, pad = 20, color = 'white');
```

```
boxplot_correlation(df, 'hour', 'consumption',  
                    palette = colors)
```



```
daytimes = ['midnight', 'early_morning', 'late_morning', 'afternoon',
            'evening', 'night']
```

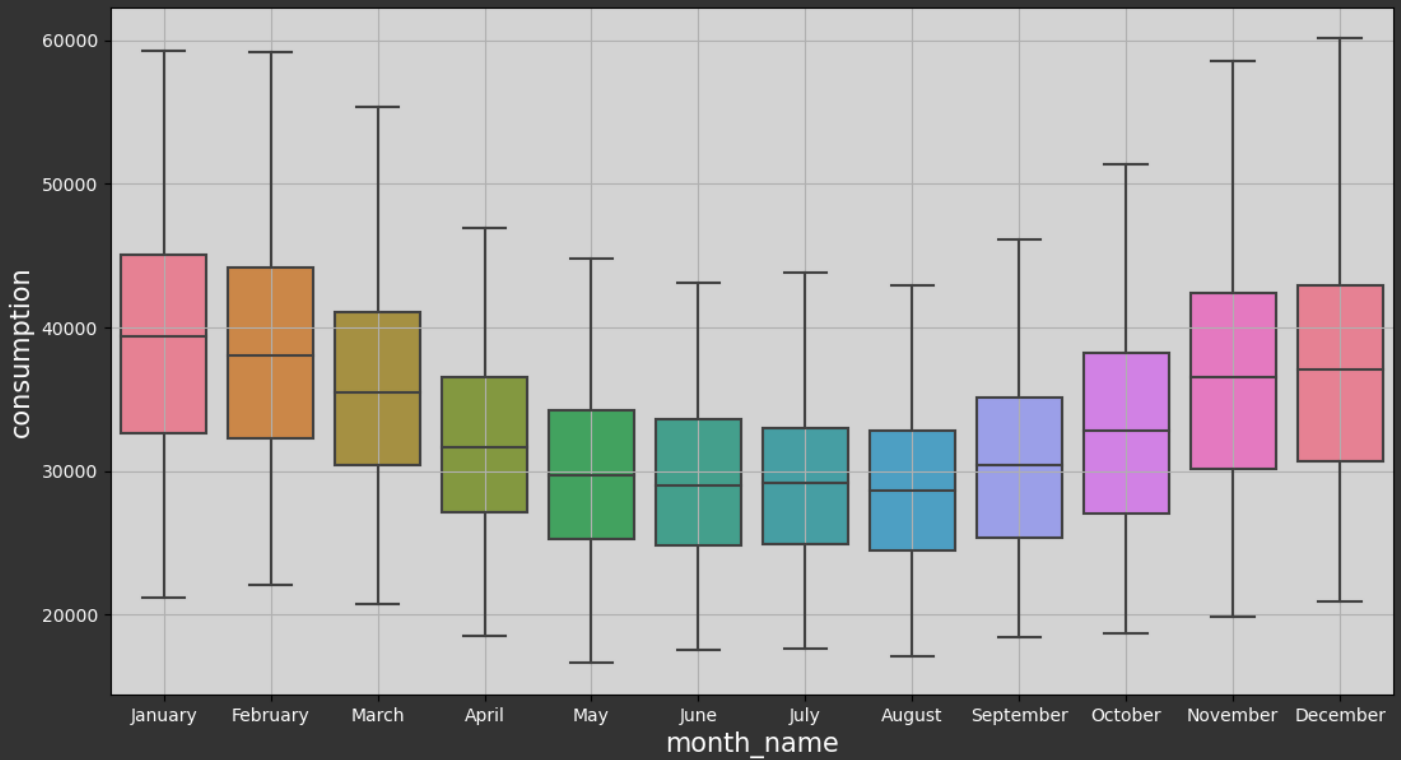
```
boxplot_correlation(df, 'time_of_day', 'consumption',
                    order = daytimes, palette = colors)
```



```
months = ['January', 'February', 'March', 'April', 'May', 'June', 'July',
          'August', 'September', 'October', 'November', 'December']
```

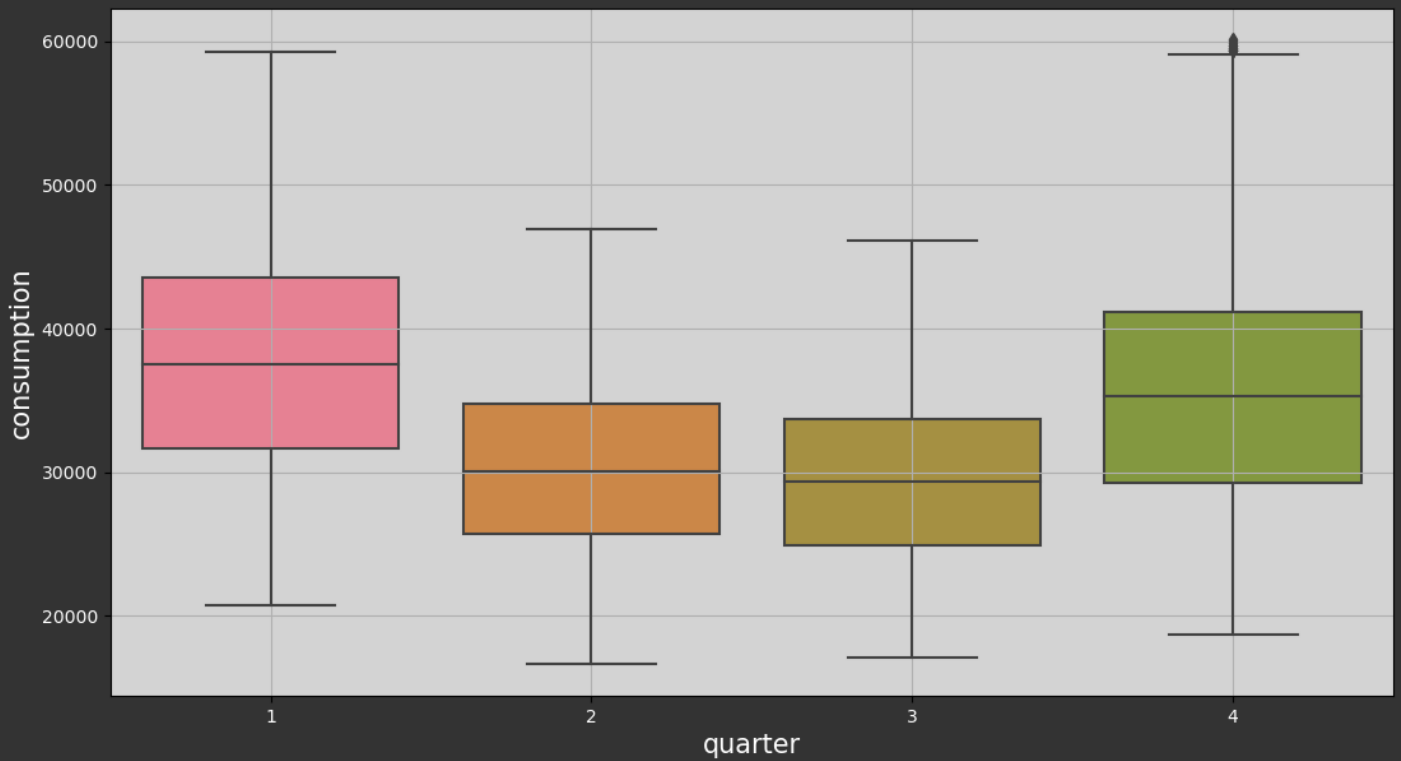
```
boxplot_correlation(df, 'month_name', 'consumption',
                    order = months, palette = colors)
```

Feature Correlation: Month\_name - Consumption



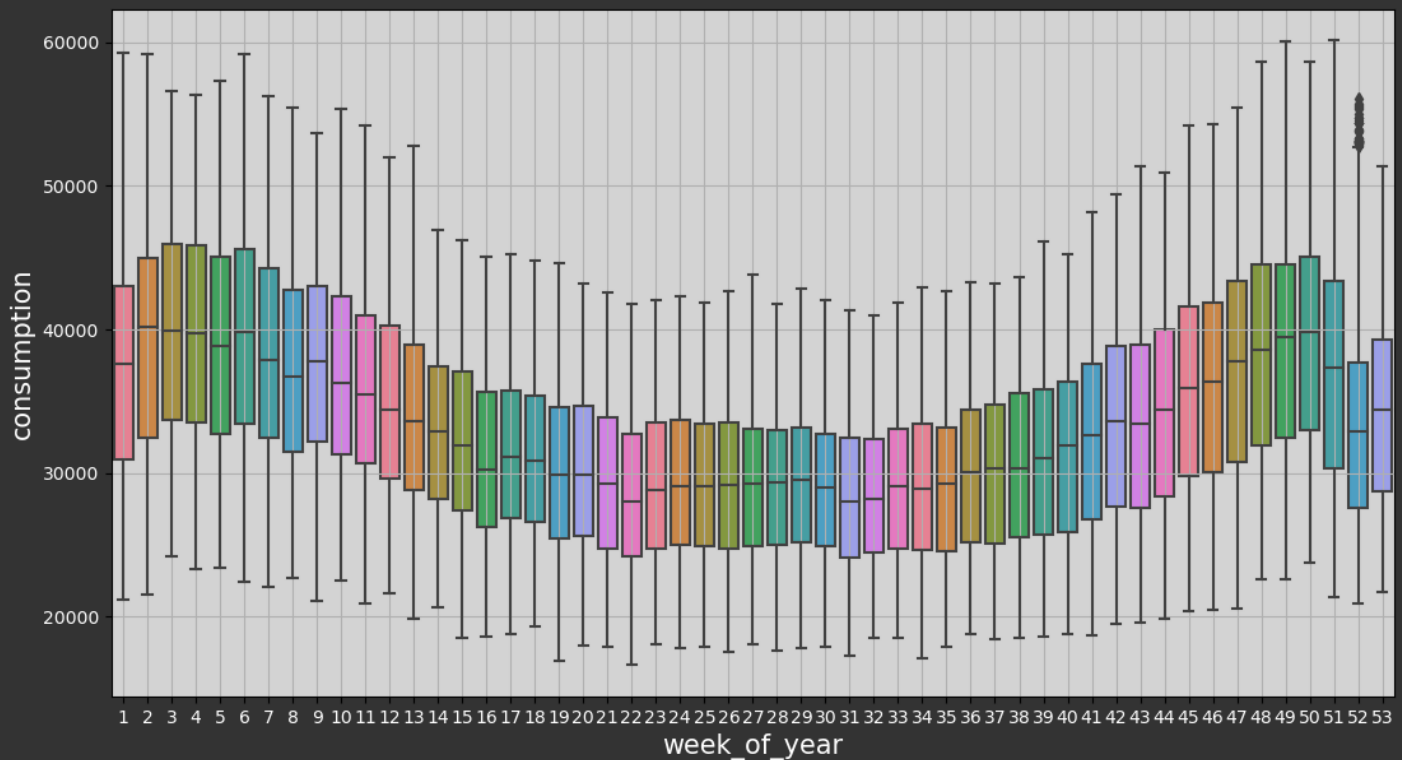
```
boxplot_correlation(df, 'quarter', 'consumption',  
                    palette = colors)
```

Feature Correlation: Quarter - Consumption



```
boxplot_correlation(df, 'week_of_year', 'consumption',  
                    palette = colors)
```

## Feature Correlation: Week\_of\_year - Consumption



### Converting week\_of\_year to float for model training

```
df['week_of_year'] = df.week_of_year.astype(float)
```

Sections: ● [Top](#) ● [The Data](#) ● [Feature Engineering](#) ● [Investigating Correlation](#) ● [Lag Features](#) ● [Splitting](#) ● [The Model](#) ● [Results with Traditional Split](#) ● [Using Cross-Validation](#) ● [Making Future Predictions](#)

## Lag Features

### Forecasting Horizon & Lag Features

- Lag Features: telling the model to look back into the past, and use the target value for that many days back as a new feature fed into the model
- in this case, it is a value from the energy usage column
- we are saving this column as a dict called target\_map
- using 364 as the increment, because it is perfectly divisible by 7 and will line up days of the week
- the forecast horizon cannot extend farther than the shortest lag chosen
- using map() to map the values

```
def year_lags(df, target_column, lag_label_list):  
  
    target_map = df[target_column].to_dict()  
    inputs = lag_label_list.copy()  
  
    for tup in inputs:
```

```
df[tup[1]] = (df.index - pd.Timedelta(tup[0])).map(target_map)

return df
```

```
lag_label_list = [('364 days', 'one_year'),
                  ('728 days', 'two_year'),
                  ('1092 days', 'three_year')]
```

```
df_lags = df.copy()
df_lags = year_lags(df_lags, 'consumption', lag_label_list)
```

```
missing_values(df_lags[['one_year', 'two_year', 'three_year']])
```

Columns and Missing Values

	column name	missing
0	one_year	18098
1	two_year	35380
2	three_year	52854

Sections: ● [Top](#) ● [The Data](#) ● [Feature Engineering](#) ● [Investigating Correlation](#) ● [Lag Features](#) ● [Splitting](#) ● [The Model](#) ● [Results with Traditional Split](#) ● [Using Cross-Validation](#) ● [Making Future Predictions](#)

## Splitting the Data

- Because I noticed that I actually got better results training without the lag features added, I will be running each operation on both the data without lag features and with lag features for comparison.

## Traditional Train-Test Split

- training data will be the first 75% of the data
- test data will be everything the last 25% of the data

```
split_point = round(len(df) * .75)
```

```
train_data = df[: split_point]
test_data = df[split_point :]
train_data = pd.get_dummies(data=train_data, columns = ['time_of_day'])
test_data = pd.get_dummies(data=test_data, columns = ['time_of_day'])
```

```
train_data_lags = df_lags[: split_point]
test_data_lags = df_lags[split_point :]
train_data_lags = pd.get_dummies(data=train_data_lags, columns = ['time_of_day'])
test_data_lags = pd.get_dummies(data=test_data_lags, columns = ['time_of_day'])
```

## Checking the point at which the training and testing divide

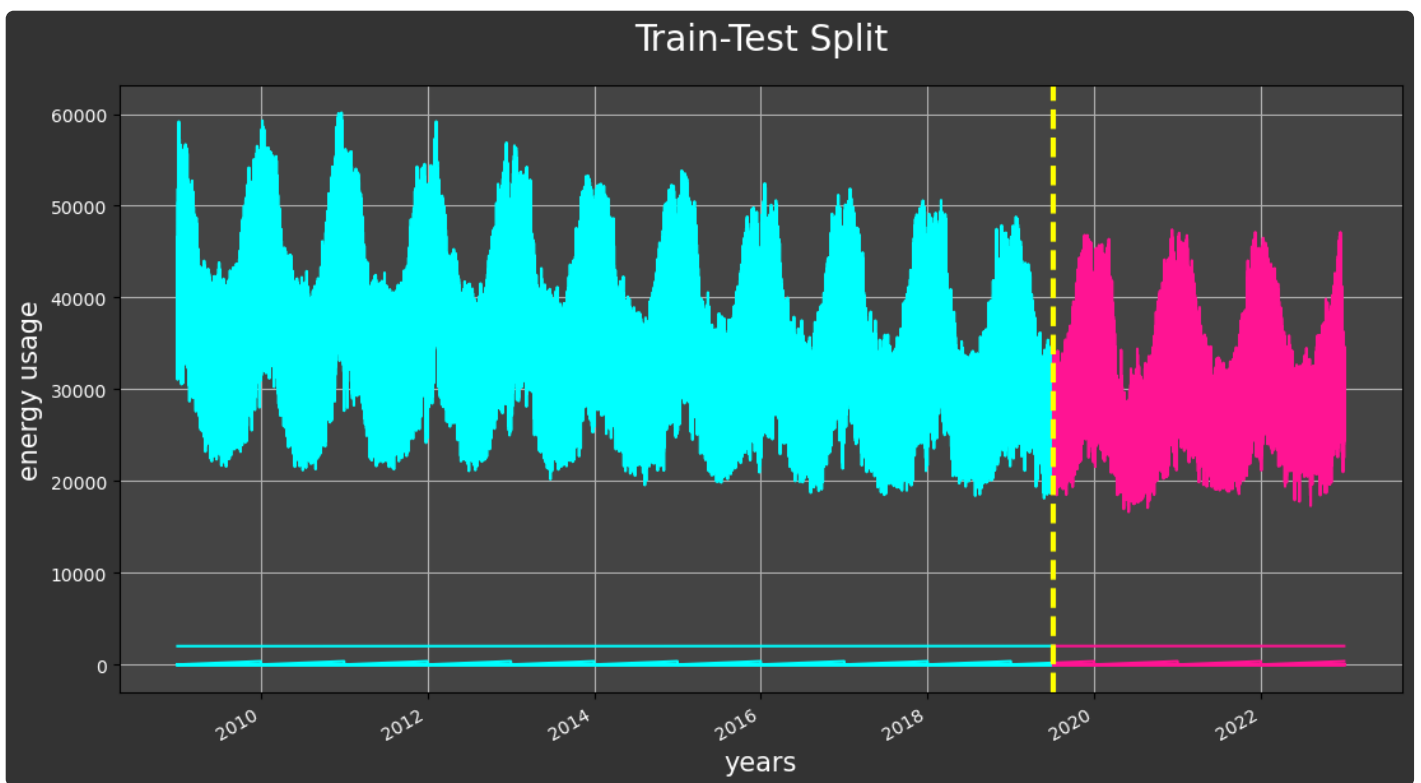
```
d(train_data.tail(1))
d(test_data.head(1))
```

	consumption	holiday	hour	weekday	weekday_name	month	month_name	quarter	year	week_of_year	date
datetime											
2019-07-06 12:00:00	27249	0	12	5	Saturday	7	July	3	2019	27	

	consumption	holiday	hour	weekday	weekday_name	month	month_name	quarter	year	week_of_year	date
datetime											
2019-07-06 12:30:00	27055	0	12	5	Saturday	7	July	3	2019	27	

## Visualizing the Traditional Train-Test Split

```
fig, ax = plt.subplots(figsize = (13, 7), facecolor = '#333333')
ax.set_facecolor('#444444')
train_data.plot(ax = ax, label = 'Training Data', color = 'cyan');
test_data.plot(ax = ax, label = 'Testing Data', color = 'deeppink');
ax.grid()
ax.axvline('2019-07-06', color = 'yellow', ls = '--', linewidth = 3)
plt.xlabel('years', color = 'white', fontsize = 15)
plt.ylabel('energy usage', color = 'white', fontsize = 15)
plt.xticks(color='white'); plt.yticks(color='white');
plt.title('Train-Test Split', fontsize = 20, pad = 20, color = 'white');
plt.legend().remove()
```



## Defining Features & Targets

```
features = ['hour', 'weekday', 'month', 'quarter', 'year',  
            'week_of_year', 'day_of_year', 'time_of_day_afternoon',  
            'time_of_day_early_morning', 'time_of_day_evening',  
            'time_of_day_late_morning', 'time_of_day_midnight',  
            'time_of_day_night']  
  
features_lags = ['hour', 'weekday', 'month', 'quarter', 'year',  
                'week_of_year', 'day_of_year', 'time_of_day_afternoon',  
                'time_of_day_early_morning', 'time_of_day_evening',  
                'time_of_day_late_morning', 'time_of_day_midnight',  
                'time_of_day_night', 'one_year', 'two_year', 'three_year']  
  
target = "consumption"
```

## Defining Training & Testing Data with Features

```
train_in = train_data[features]  
train_out = train_data[target]  
test_in = test_data[features]  
test_out = test_data[target]
```

```
train_in_lags = train_data_lags[features_lags]  
train_out_lags = train_data_lags[target]  
test_in_lags = test_data_lags[features_lags]  
test_out_lags = test_data_lags[target]
```

## The XGBoost Regressor Model

### Creating the Model (without lags)

```
xgb_regressor = XGBRegressor(n_estimators = 1200,  
                             learning_rate = 0.012,  
                             early_stopping_rounds = 23)
```

### Fitting the Model (without lags)

```
xgb_regressor.fit(train_in, train_out,  
                  eval_set = [(train_in, train_out),  
                              (test_in, test_out)], verbose = 25)
```

```
[0] validation_0-rmse:35068.46435    validation_1-rmse:29587.12720  
[25]  validation_0-rmse:26018.01179    validation_1-rmse:21235.48885  
[50]  validation_0-rmse:19337.02538    validation_1-rmse:15148.51055  
[75]  validation_0-rmse:14411.06888    validation_1-rmse:10692.43309  
[100] validation_0-rmse:10787.75506    validation_1-rmse:7519.52345  
[125] validation_0-rmse:8131.03112     validation_1-rmse:5323.82623  
[150] validation_0-rmse:6192.47195     validation_1-rmse:3833.85620  
[175] validation_0-rmse:4791.75784     validation_1-rmse:2961.30137  
[200] validation_0-rmse:3789.02170     validation_1-rmse:2547.03297  
[225] validation_0-rmse:3086.91956     validation_1-rmse:2420.27304  
[250] validation_0-rmse:2602.61653     validation_1-rmse:2429.93979  
[257] validation_0-rmse:2496.84033     validation_1-rmse:2442.06684
```

```
XGBRegressor(base_score=None, booster=None, callbacks=None,  
             colsample_bylevel=None, colsample_bynode=None,  
             colsample_bytree=None, early_stopping_rounds=23,  
             enable_categorical=False, eval_metric=None, feature_types=None,  
             gamma=None, gpu_id=None, grow_policy=None, importance_type=None,  
             interaction_constraints=None, learning_rate=0.012, max_bin=None,  
             max_cat_threshold=None, max_cat_to_onehot=None,  
             max_delta_step=None, max_depth=None, max_leaves=None,  
             min_child_weight=None, missing=nan, monotone_constraints=None,  
             n_estimators=1200, n_jobs=None, num_parallel_tree=None,  
             predictor=None, random_state=None, ...)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.  
XGBRegressor



```
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=None, early_stopping_rounds=23,
             enable_categorical=False, eval_metric=None, feature_types=None,
             gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
             interaction_constraints=None, learning_rate=0.012, max_bin=None,
             max_cat_threshold=None, max_cat_to_onehot=None,
             max_delta_step=None, max_depth=None, max_leaves=None,
             min_child_weight=None, missing=nan, monotone_constraints=None,
             n_estimators=1200, n_jobs=None, num_parallel_tree=None,
             predictor=None, random_state=None, ...)
```

## Fitting the Model (with lags)

```
xgb_regressor_lags = XGBRegressor(n_estimators = 1200,
                                   learning_rate = 0.012,
                                   early_stopping_rounds = 23)
```

```
xgb_regressor_lags.fit(train_in_lags, train_out_lags,
                       eval_set = [(train_in_lags, train_out_lags),
                                   (test_in_lags, test_out_lags)],
                       verbose = 25)
```

```
[0] validation_0-rmse:35067.37778    validation_1-rmse:29621.32089
[25]  validation_0-rmse:25992.89494    validation_1-rmse:21982.27167
[50]  validation_0-rmse:19296.64994    validation_1-rmse:16362.81470
[75]  validation_0-rmse:14362.19803    validation_1-rmse:12236.19017
[100] validation_0-rmse:10732.97496    validation_1-rmse:9220.63188
[125] validation_0-rmse:8075.19740     validation_1-rmse:7039.05742
[150] validation_0-rmse:6141.15746     validation_1-rmse:5477.98424
[175] validation_0-rmse:4744.45133     validation_1-rmse:4392.53113
[200] validation_0-rmse:3752.23105     validation_1-rmse:3673.74192
[225] validation_0-rmse:3063.21492     validation_1-rmse:3212.91990
[250] validation_0-rmse:2594.90743     validation_1-rmse:2925.86404
[275] validation_0-rmse:2282.94267     validation_1-rmse:2750.60546
[300] validation_0-rmse:2073.32089     validation_1-rmse:2662.13021
[325] validation_0-rmse:1937.22774     validation_1-rmse:2611.90552
[350] validation_0-rmse:1836.60203     validation_1-rmse:2585.33120
[375] validation_0-rmse:1762.79711     validation_1-rmse:2569.20642
[400] validation_0-rmse:1706.81078     validation_1-rmse:2561.84587
[425] validation_0-rmse:1670.78904     validation_1-rmse:2555.93349
[450] validation_0-rmse:1644.08868     validation_1-rmse:2551.10259
[475] validation_0-rmse:1621.20190     validation_1-rmse:2548.89837
[500] validation_0-rmse:1604.13354     validation_1-rmse:2547.76779
```

```
[525] validation_0-rmse:1587.04879 validation_1-rmse:2549.01441
[531] validation_0-rmse:1583.52664 validation_1-rmse:2549.17672
```

```
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=None, early_stopping_rounds=23,
             enable_categorical=False, eval_metric=None, feature_types=None,
             gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
             interaction_constraints=None, learning_rate=0.012, max_bin=None,
             max_cat_threshold=None, max_cat_to_onehot=None,
             max_delta_step=None, max_depth=None, max_leaves=None,
             min_child_weight=None, missing=nan, monotone_constraints=None,
             n_estimators=1200, n_jobs=None, num_parallel_tree=None,
             predictor=None, random_state=None, ...)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with [nbviewer.org](https://nbviewer.org).  
XGBRegressor

```
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=None, early_stopping_rounds=23,
             enable_categorical=False, eval_metric=None, feature_types=None,
             gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
             interaction_constraints=None, learning_rate=0.012, max_bin=None,
             max_cat_threshold=None, max_cat_to_onehot=None,
             max_delta_step=None, max_depth=None, max_leaves=None,
             min_child_weight=None, missing=nan, monotone_constraints=None,
             n_estimators=1200, n_jobs=None, num_parallel_tree=None,
             predictor=None, random_state=None, ...)
```

## Feature Importance (without lags)

```
feature_importance = pd.DataFrame(data = xgb_regressor.feature_importances_,
                                  index = xgb_regressor.feature_names_in_,
                                  columns = ['importance'])

feature_importance = feature_importance[feature_importance.importance > 0.0005]

multi([(feature_importance.sort_values('importance', ascending = False),
        'Feature Importance (no lags)'),
      (feature_importance_lags.sort_values('importance', ascending = False),
        'Feature Importance (with lags)')], precision = 4)
```

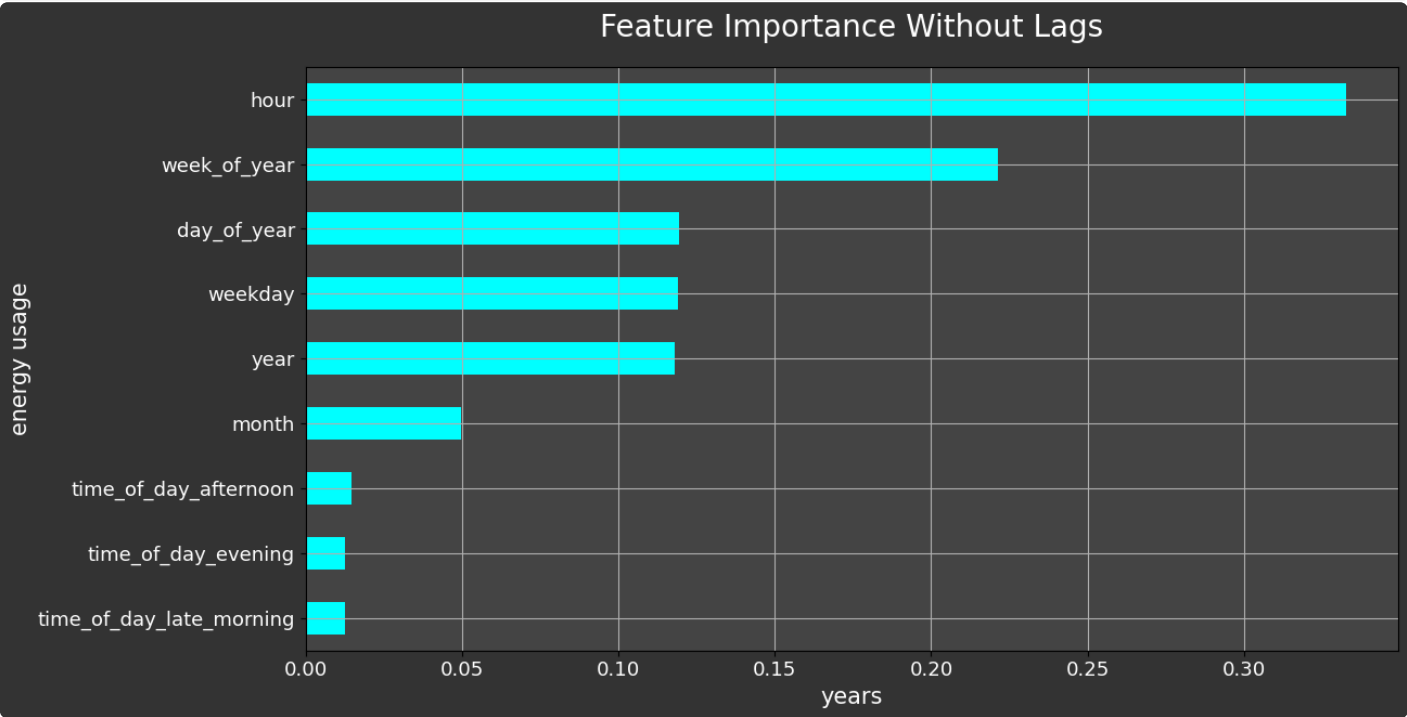
Feature Importance (no lags)		Feature Importance (with lags)	
	importance		importance
hour	0.3328	hour	0.3328
week_of_year	0.2212	week_of_year	0.2212
day_of_year	0.1193	day_of_year	0.1193
weekday	0.1191	weekday	0.1191
year	0.1179	year	0.1179
month	0.0498	month	0.0498
time_of_day_afternoon	0.0147	time_of_day_afternoon	0.0147
time_of_day_evening	0.0127	time_of_day_evening	0.0127
time_of_day_late_morning	0.0125	time_of_day_late_morning	0.0125

Plotting Feature Importance

```
fig, ax = plt.subplots(figsize = (13, 7), facecolor = '#333333')
ax.set_facecolor('#444444')

feature_importance.sort_values('importance').plot(kind = 'barh',
                                                    color = 'cyan',
                                                    ax = ax)

ax.grid()
plt.xlabel('years', color = 'white', fontsize = 15)
plt.ylabel('energy usage', color = 'white', fontsize = 15)
plt.xticks(color='white', size = 13); plt.yticks(color='white', size = 13);
plt.title('Feature Importance Without Lags', fontsize = 20, pad = 20, color = 'white');
plt.legend().remove()
```



## Results with Traditional Train-Test Split

- Because there is no benefit to having the lag data added, for the remainder of the project, I will be working with only the data without lags added.

## Predicting with Test Data


```
test_data['prediction'] = xgb_regressor.predict(test_in)
```

## Merging Predictions into Original DF

```
df = df.merge(test_data['prediction'],
              how = 'left',
              left_index = True,
              right_index = True)
```

## Plotting Predictions vs Original Data

```
plt.figure(facecolor = '#333333');
ax = df['consumption'].plot(figsize = (13, 7), color = 'cyan');
ax.set_facecolor('#444444');
df['prediction'].plot(ax = ax, style = '.', color = 'deeppink');
plt.legend();
ax.grid()
plt.xlabel('years', color = 'white', fontsize = 15)
plt.ylabel('energy usage', color = 'white', fontsize = 15)
plt.xticks(color='white', size = 13); plt.yticks(color='white', size = 13);
plt.title('Train-Test Split', fontsize = 20, pad = 20, color = 'white');
plt.legend().remove()
ax.set_title('Actual Data vs Predicted', color = 'white', size = 20, pad = 20);
```



### Predictions vs Targets: One Week Excerpt

```
plt.figure(figsize = (13, 7), facecolor = '#333333')
ax = df.loc[(df.index > '2019-07-06') & (df.index <= '2019-07-13')]\
    ['consumption'].plot(color = 'cyan', linewidth = 3);
df.loc[(df.index > '2019-07-06') & (df.index <= '2019-07-13')]\
    ['prediction'].plot(style = '.', color = 'yellow', markersize = 9);
ax.set_facecolor('#444444')
plt.legend();
ax.grid()
plt.xlabel('days', color = 'white', fontsize = 15)
plt.ylabel('energy usage', color = 'white', fontsize = 15)
plt.tick_params(labelcolor = 'white', which = 'both')
plt.title('Training Data vs Predictions: 1 Week Excerpt', fontsize = 20, pad = 20, color = 'white')
plt.legend(fontsize = 15, labelcolor = 'white', facecolor = '#333333');
```



```

        'Average Sharpe Ratio'],
        name = "Resulting Metrics")

see(results, 'Overall Metrics for Accuracy and Model Performance')

return df

```

```
pred_actual = df.loc[df.prediction.notna()]
```

```

pred_actual = get_accuracy(pred_actual, 'consumption', 'prediction')

head_tail_horz(pred_actual[['consumption', 'prediction', 'abs_acc',
                             'rel_acc', 'sharpe']], 5, 'Predictions vs Actual Targets')

```

### Overall Metrics for Accuracy and Model Performance

Resulting Metrics	
Average RMSE	2,414.50
Average Absolute Accuracy	93.58%
Average Relative Accuracy	91.91%
Average Sharpe Ratio	0.33

### Predictions vs Actual Targets

head(5)

	consumption	prediction	abs_acc	rel_acc	sharpe
datetime					
2019-07-06	27,055	24,169.05	88.06	87.51	0.51
2019-07-06	26,790	24,147.77	89.06	88.56	0.47
2019-07-06	26,434	24,147.77	90.53	90.10	0.40
2019-07-06	26,204	24,147.77	91.48	91.10	0.36
2019-07-06	25,862	24,147.77	92.90	92.58	0.30

tail(5)

	consumption	prediction	abs_acc	rel_acc	sharpe
datetime					
2022-12-31	25,634	32,205.60	79.59	71.56	1.16
2022-12-31	24,788	29,192.30	84.91	80.94	0.78
2022-12-31	24,365	29,192.30	83.46	79.11	0.85
2022-12-31	24,766	26,653.43	92.92	91.83	0.33
2022-12-31	24,843	26,653.43	93.21	92.16	0.32


## Plotting the Sharpe Ratio

```

plt.figure(figsize = (13, 7), facecolor = '#333333')
ax = plt.axes()
pred_actual.sharpe.plot(kind="hist",
                        bins=100,
                        color = 'cyan',
                        edgecolor = 'white',
                        ax = ax)

ax.set_facecolor('#444444')
plt.legend();
ax.grid()
plt.xlabel('sharpe ratio', color = 'white', fontsize = 15)
plt.ylabel('count', color = 'white', fontsize = 15)
plt.tick_params(labelcolor = 'white', which = 'both')
plt.title('Overall Distribution of Sharpe Ratio Metrics',
          fontsize = 20, pad = 20, color = 'white');
plt.legend().remove()

```



## Daily Error - Investigating the Days with the Highest and Lowest Error

```

def get_daily_error(df, actual_col, pred_col, num_examples,
                    ascending = False):

    temp = df[[actual_col, pred_col]].copy()
    temp['date'] = temp.index.strftime('%A, %b %d, %Y')
    temp['error'] = np.abs(df[actual_col] - df[pred_col])

    results = temp.sort_values("error", ascending = ascending)

    error_style = {'error': [{'selector': '',

```



```

        'props': [('color', 'red'),
                  ('font-weight', 'bold'),
                  ('padding-right', '15px'),
                  ('padding-left', '15px')]]],
    'date': [{'selector': 'td',
              'props': [('color', 'blue'),
                        ('font-weight', 'bold'),
                        ('padding-right', '15px'),
                        ('padding-left', '15px')]]],
    'prediction': [{'selector': 'td',
                   'props': [('padding-right', '25px'),
                             ('padding-left', '15px')]]}]

if ascending == True:
    pretty(f'Daily error for the {num_examples} days with the lowest error:',
          fontsize = 4)
else:
    pretty(f'Daily error for the {num_examples} days with the highest error:',
          fontsize = 4)

return results[['date',
                'error',
                pred_col,
                actual_col]].head(num_examples).style.hide(axis='index')\
    .set_table_styles(error_style)\
    .format(precision=3, thousands=","))

```

```
get_daily_error(pred_actual, 'consumption', 'prediction', 10)
```

Daily error for the 10 days with the highest error:

	date	error	prediction	consumption
	Monday, Jan 03, 2022	12,227.672	38,637.672	26,410
	Sunday, Dec 25, 2022	11,386.129	34,748.129	23,362
	Monday, May 25, 2020	11,337.658	29,313.658	17,976
	Monday, May 25, 2020	11,248.658	29,313.658	18,065
	Monday, Jan 03, 2022	11,133.672	38,637.672	27,504
	Sunday, Dec 25, 2022	11,090.129	34,748.129	23,658
	Monday, May 25, 2020	10,898.234	29,617.234	18,719
	Sunday, Dec 25, 2022	10,891.258	34,945.258	24,054
	Monday, Apr 13, 2020	10,769.680	30,451.680	19,682
	Monday, Apr 13, 2020	10,671.783	30,159.783	19,488

```
get_daily_error(pred_actual, 'consumption', 'prediction', 10,
                ascending = True)
```

Daily error for the 10 days with the lowest error:

	date	error	prediction	consumption
--	------	-------	------------	-------------

	date	error	prediction	consumption
	Monday, Oct 14, 2019	0.107	25,882.107	25,882
	Saturday, Dec 14, 2019	0.250	26,106.250	26,106
	Wednesday, Oct 13, 2021	0.395	32,784.605	32,785
	Tuesday, Feb 09, 2021	0.395	40,235.605	40,236
	Tuesday, Nov 26, 2019	0.443	26,390.443	26,390
	Friday, Aug 27, 2021	0.479	28,798.521	28,799
	Saturday, Oct 30, 2021	0.508	23,858.508	23,858
	Thursday, Dec 02, 2021	0.557	31,655.443	31,656
	Sunday, Jun 27, 2021	0.600	20,580.600	20,580
	Monday, Mar 14, 2022	0.686	24,652.314	24,653

---

Sections: ● [Top](#) ● [The Data](#) ● [Feature Engineering](#) ● [Investigating Correlation](#) ● [Lag Features](#) ● [Splitting](#) ● [The Model](#) ● [Results with Traditional Split](#) ● [Using Cross-Validation](#) ● [Making Future Predictions](#)

---

## Training the Model Using Cross Validation

### What is Cross Validation?

- using sklearn TimeSeriesSplit()
- test size is 1 year of hourly records
- a gap of 24 puts a 1 day gap between the end of a training set and beginning of a test set
- must be sure time series data is sorted so that the time series split will work

```
time_series_split = TimeSeriesSplit(n_splits = 5,
                                    test_size = (24 * 365 * 1),
                                    gap = 24)

df = df.sort_index()
```

### This creates a time series split generator object

- this object will be applied across the data
- it will loop over the data for as many splits as are passed

```
time_series_split
```

```
TimeSeriesSplit(gap=24, max_train_size=None, n_splits=5, test_size=8760)
```

### Visualizing the TimeSeriesSplit() process

- for each fold, the model goes back and tests the testing data in 1 year increments, independently from the rest of the testing data

- this is a very good approach when the dataset has a large number of records

```
for train_index, validation_index in time_series_split.split(df):
    break
pretty(train_index, 'Indices of training data:')
pretty(validation_index, 'Indices of validation data:'); sp()
```

Indices of training data:

[0 1 2 ... 200849 200850 200851]

Indices of validation data:

[200876 200877 200878 ... 209633 209634 209635]

```
fig, axs = plt.subplots(5, 1, figsize = (13, 10),
                        sharex = True, facecolor = "#222222")

split = 0
for train_index, validation_index in time_series_split.split(df):
    training = df.iloc[train_index]
    validation = df.iloc[validation_index]

    ax = axs[split]
    ax.set_facecolor('#444444')
    ax.set_title(f'Training & Testing split for split no. {split}',
                fontsize = 15, color = 'white')
    ax.tick_params(color = 'white', labelcolor = 'white')

    training['consumption'].plot(ax = ax,
                                label = 'Training Set',
                                color = 'cyan')
    validation['consumption'].plot(ax = ax,
                                   label = 'Validation Set',
                                   color = 'deeppink')

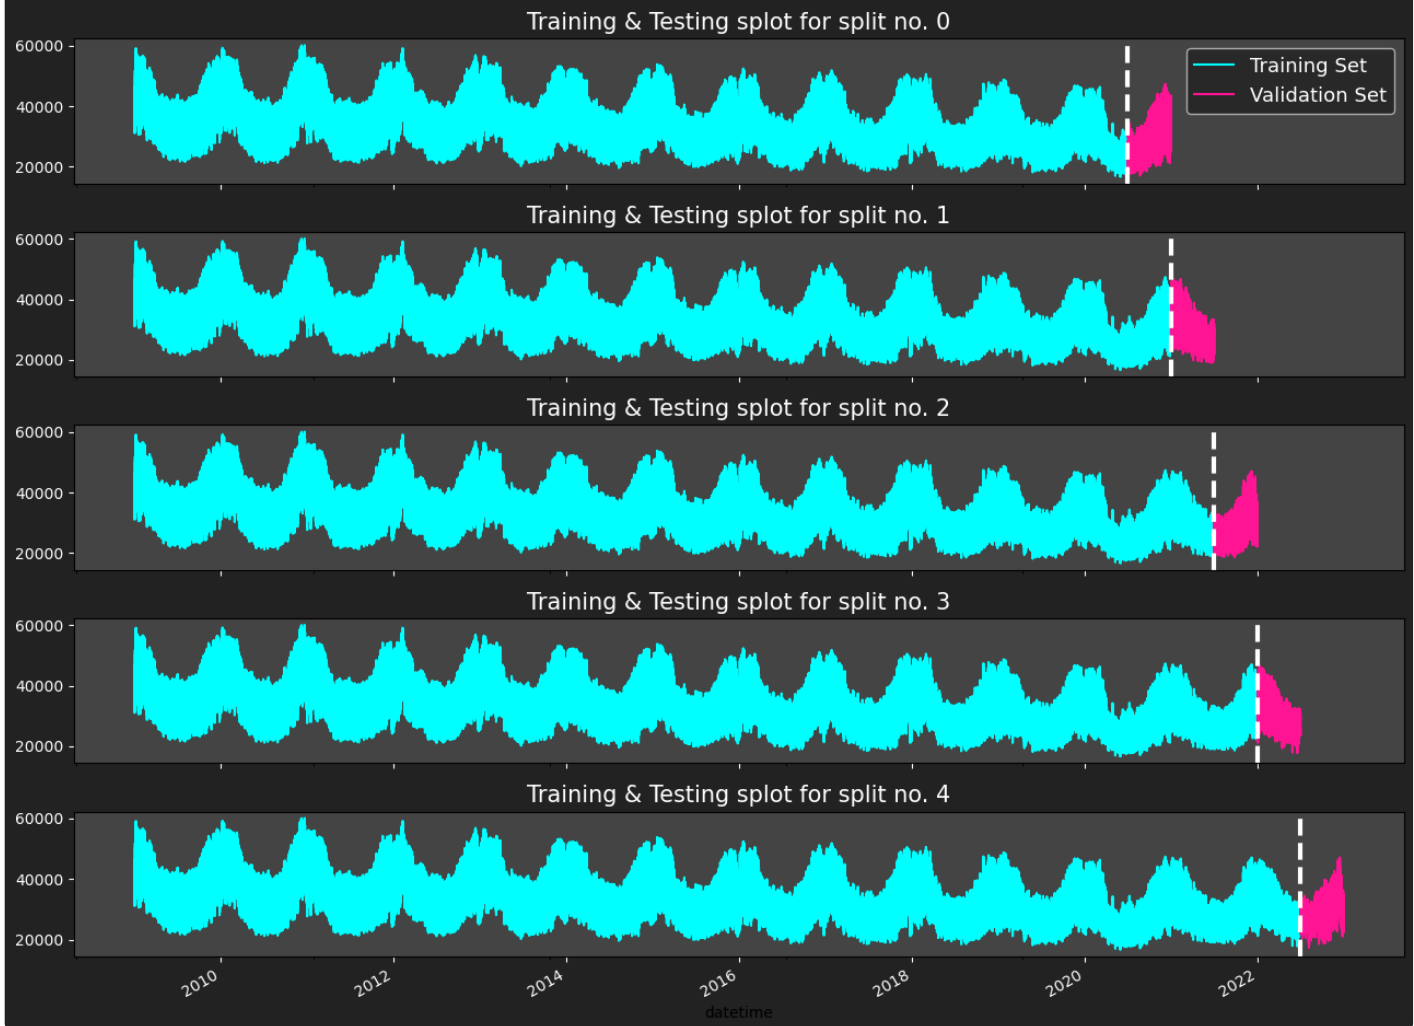
    axs[split].axvline(validation.index.min(),
                      color = 'white', ls = "--", linewidth = 3)

    axs[0].legend(fontsize=13, facecolor = '#222222', labelcolor = 'white')

    split += 1

plt.suptitle('5 Different Folds of Cross Validation', color = 'white', fontsize = 24)
plt.tight_layout()
```

## 5 Different Folds of Cross Validation



## Training with Cross Validation

- this will repeat the steps above in order to now train with cross validation and the `TimeSeriesSplit()` method

```
def cross_validation_train(df, target_column,
                          lags = False, lag_label_list = None,
                          splits = 5, test_size = (24 * 365 * 1),
                          gap = 24, base_score = 0.5, booster = 'gbtree',
                          n_estimators = 1200, early_stopping_rounds = 50,
                          objective = 'reg:squarederror', max_depth = 3,
                          learning_rate = 0.01, verbose = 25):

    from sklearn.metrics import mean_squared_error
    from sklearn.model_selection import TimeSeriesSplit

    # ..... Establishing Features & Targets ..... #

    df = featurize_datetime_index(df)
    target = target_column

    # ..... Lag Features & Feature Cleanup ..... #
```

```

if lags:
    df = year_lags(df, target_column, lag_label_list)
    features = list(df.columns)

else:
    features = list(df.columns)

features.remove('weekday_name')
features.remove('month_name')
features.remove('time_of_day')
features.remove(target)

# ..... Correcting Datatypes for Model ..... #

df['week_of_year'] = df.week_of_year.astype(float)

# ..... Setting Up Cross-Validation ..... #

split = 1
prediction_log = []
target_log = []
rmse_log = []

df = df.sort_index()
time_series_split = TimeSeriesSplit(n_splits = splits,
                                     test_size = test_size,
                                     gap = gap)

# ..... Running Cross-Validation Training ..... #

for train_index, validation_index in time_series_split.split(df):
    pretty(f'Training split: {split} of {splits}')
    training = df.iloc[train_index]
    validation = df.iloc[validation_index]

    train_in = training[features]
    train_out = training[target]
    test_in = validation[features]
    test_out = validation[target]

    model = XGBRegressor(base_score = base_score,
                         booster = booster,
                         n_estimators = n_estimators,
                         early_stopping_rounds = early_stopping_rounds,
                         objective = objective,
                         max_depth = max_depth,
                         learning_rate = learning_rate)

    model.fit(train_in, train_out,
              eval_set = [(train_in, train_out),
                          (test_in, test_out)],

```

```

        verbose = verbose)

    targets = list(test_out)
    target_log += targets

    prediction = model.predict(test_in)
    prediction_log += list(prediction)

    rmse = np.sqrt(mean_squared_error(test_out, prediction))
    rmse_log.append((f'split {split}', rmse))

    split += 1

# ..... Compiling Training & Testing Results ..... #

results = pd.DataFrame(pd.concat([pd.Series(prediction_log),
                                   pd.Series(target_log)],
                                   axis = 1))
results.columns = ['prediction', 'actual']

pred_actual = get_accuracy(results, 'actual', 'prediction')

print('')

head_tail_horz(pred_actual[['actual', 'prediction', 'abs_acc',
                             'rel_acc', 'sharpe']], 5, 'Predictions vs Actual Targets')

return pred_actual, rmse_log

```

## Cross-Validation Results without Lag Features

```

cross_results, rmses = cross_validation_train(df,
                                              target_column = "consumption",
                                              learning_rate = 0.012,
                                              verbose = 500,
                                              splits = 16)

```

### Training split: 1 of 16

```

[0] validation_0-rmse:36927.43252 validation_1-rmse:34875.18364
[500] validation_0-rmse:2056.87214 validation_1-rmse:2423.05766
[1000] validation_0-rmse:1676.66795 validation_1-rmse:2181.13980
[1199] validation_0-rmse:1616.77561 validation_1-rmse:2155.05839

```

### Training split: 2 of 16

```

[0] validation_0-rmse:36775.76007 validation_1-rmse:32375.93347
[265] validation_0-rmse:3149.83870 validation_1-rmse:3050.08925

```

### Training split: 3 of 16

```

[0] validation_0-rmse:36479.99831 validation_1-rmse:33557.74521
[298] validation_0-rmse:2863.89178 validation_1-rmse:2735.91368

```

#### Training split: 4 of 16

```
[0] validation_0-rmse:36295.33002 validation_1-rmse:32125.01992
[500] validation_0-rmse:2137.87922 validation_1-rmse:2715.88639
[1000] validation_0-rmse:1753.45240 validation_1-rmse:2321.19621
[1199] validation_0-rmse:1693.75240 validation_1-rmse:2266.00976
```

#### Training split: 5 of 16

```
[0] validation_0-rmse:36049.16304 validation_1-rmse:32516.82046
[281] validation_0-rmse:3059.10355 validation_1-rmse:3166.45254
```

#### Training split: 6 of 16

```
[0] validation_0-rmse:35853.96617 validation_1-rmse:31558.53324
[500] validation_0-rmse:2192.45384 validation_1-rmse:2735.75052
[1000] validation_0-rmse:1807.99360 validation_1-rmse:2231.06479
[1199] validation_0-rmse:1746.01944 validation_1-rmse:2185.40512
```

#### Training split: 7 of 16

```
[0] validation_0-rmse:35629.26492 validation_1-rmse:32723.94519
[500] validation_0-rmse:2209.88738 validation_1-rmse:2666.12457
[803] validation_0-rmse:1906.31098 validation_1-rmse:2620.93845
```

#### Training split: 8 of 16

```
[0] validation_0-rmse:35484.56719 validation_1-rmse:30590.92882
[500] validation_0-rmse:2238.18373 validation_1-rmse:2665.97085
[1000] validation_0-rmse:1874.88304 validation_1-rmse:2380.88144
[1199] validation_0-rmse:1815.11051 validation_1-rmse:2338.74912
```

#### Training split: 9 of 16

```
[0] validation_0-rmse:35256.84386 validation_1-rmse:31274.43857
[291] validation_0-rmse:3042.73232 validation_1-rmse:3021.40147
```

#### Training split: 10 of 16

```
[0] validation_0-rmse:35080.13209 validation_1-rmse:30334.97187
[500] validation_0-rmse:2293.31344 validation_1-rmse:2368.77382
[1000] validation_0-rmse:1929.70739 validation_1-rmse:1904.35239
[1199] validation_0-rmse:1875.09340 validation_1-rmse:1853.25663
```

#### Training split: 11 of 16

```
[0] validation_0-rmse:34879.05495 validation_1-rmse:28467.95624
[214] validation_0-rmse:4217.73427 validation_1-rmse:3978.68606
```

#### Training split: 12 of 16

```
[0] validation_0-rmse:34628.76251 validation_1-rmse:28969.99048
[500] validation_0-rmse:2316.41624 validation_1-rmse:2452.28254
[828] validation_0-rmse:2030.48465 validation_1-rmse:2354.46986
```

#### Training split: 13 of 16

```
[0] validation_0-rmse:34410.39062    validation_1-rmse:30310.69798
[407] validation_0-rmse:2547.65380    validation_1-rmse:2682.01408
```

Training split: 14 of 16

```
[0] validation_0-rmse:34257.96348 validation_1-rmse:29340.27107
[500] validation_0-rmse:2362.09599 validation_1-rmse:1989.56895
[607] validation_0-rmse:2223.56214 validation_1-rmse:2001.37904
```

Training split: 15 of 16

```
[0] validation_0-rmse:34081.45769    validation_1-rmse:30150.75685
[460] validation_0-rmse:2437.94536    validation_1-rmse:2518.15912
```

Training split: 16 of 16

```
[0] validation_0-rmse:33944.73814    validation_1-rmse:29491.10969
[382] validation_0-rmse:2668.21192    validation_1-rmse:2609.05645
```

## Overall Metrics for Accuracy and Model Performance

	Resulting Metrics
Average RMSE	2,582.19
Average Absolute Accuracy	93.57%
Average Relative Accuracy	93.70%
Average Sharpe Ratio	0.33

### Predictions vs Actual Targets

	head(5)						tail(5)				
	actual	prediction	abs_acc	rel_acc	sharpe		actual	prediction	abs_acc	rel_acc	sharpe
0	41,297	44,150.49	93.54	90.99	0.46	140155	25,634	32,433.49	79.04	78.52	1.11
1	44,813	44,150.49	98.50	97.91	0.11	140156	24,788	29,763.41	83.28	84.29	0.81
2	46,139	45,633.64	98.89	98.40	0.08	140157	24,365	29,763.41	81.86	82.95	0.88
3	45,716	45,633.64	99.82	99.74	0.01	140158	24,766	27,849.91	88.93	90.26	0.50
4	45,342	45,633.64	99.36	99.08	0.05	140159	24,843	27,849.91	89.20	90.50	0.49

## Cross-Validation Results with Lag Features

[illegible]



#### Training split: 1 of 16

[0] validation\_0-rmse:36925.12428 validation\_1-rmse:34888.08778  
[500] validation\_0-rmse:2037.49828 validation\_1-rmse:2026.48655  
[576] validation\_0-rmse:1964.88895 validation\_1-rmse:2029.35903

#### Training split: 2 of 16

[0] validation\_0-rmse:36773.31063 validation\_1-rmse:32394.96397  
[335] validation\_0-rmse:2386.18647 validation\_1-rmse:1923.99288

#### Training split: 3 of 16

[0] validation\_0-rmse:36477.26414 validation\_1-rmse:33584.40726  
[444] validation\_0-rmse:2104.36729 validation\_1-rmse:2129.39891

#### Training split: 4 of 16

[0] validation\_0-rmse:36292.39954 validation\_1-rmse:32157.17855  
[500] validation\_0-rmse:2046.95073 validation\_1-rmse:2450.85211  
[902] validation\_0-rmse:1815.59204 validation\_1-rmse:2399.96702

#### Training split: 5 of 16

[0] validation\_0-rmse:36046.01723 validation\_1-rmse:32549.80008  
[500] validation\_0-rmse:2064.69275 validation\_1-rmse:2385.32174  
[1000] validation\_0-rmse:1794.49164 validation\_1-rmse:2315.57542  
[1194] validation\_0-rmse:1740.22213 validation\_1-rmse:2305.69089

#### Training split: 6 of 16

[0] validation\_0-rmse:35850.48363 validation\_1-rmse:31592.50643  
[500] validation\_0-rmse:2095.36822 validation\_1-rmse:2027.17829  
[504] validation\_0-rmse:2090.82782 validation\_1-rmse:2024.67160

#### Training split: 7 of 16

[0] validation\_0-rmse:35625.58122 validation\_1-rmse:32765.45789  
[500] validation\_0-rmse:2089.91556 validation\_1-rmse:2825.18626  
[548] validation\_0-rmse:2048.50601 validation\_1-rmse:2831.43948

#### Training split: 8 of 16

[0] validation\_0-rmse:35480.78265 validation\_1-rmse:30627.60853  
[412] validation\_0-rmse:2209.38032 validation\_1-rmse:2202.98109

#### Training split: 9 of 16

[0] validation\_0-rmse:35252.84829 validation\_1-rmse:31312.27241  
[369] validation\_0-rmse:2293.48245 validation\_1-rmse:2485.35647

#### Training split: 10 of 16

[0] validation\_0-rmse:35075.91026 validation\_1-rmse:30377.44317  
[500] validation\_0-rmse:2132.74711 validation\_1-rmse:1835.06696  
[748] validation\_0-rmse:1989.27927 validation\_1-rmse:1819.15242

#### Training split: 11 of 16

[0] validation\_0-rmse:34874.70089 validation\_1-rmse:28512.73610  
[261] validation\_0-rmse:2895.43167 validation\_1-rmse:2931.84616

Training split: 12 of 16

[0] validation\_0-rmse:34623.79921 validation\_1-rmse:28973.77077  
[400] validation\_0-rmse:2271.53341 validation\_1-rmse:2734.28776

Training split: 13 of 16

[0] validation\_0-rmse:34405.31749 validation\_1-rmse:30354.22325  
[400] validation\_0-rmse:2289.59348 validation\_1-rmse:3920.04580

Training split: 14 of 16

[0] validation\_0-rmse:34252.93483 validation\_1-rmse:29339.39192  
[495] validation\_0-rmse:2193.78422 validation\_1-rmse:1936.44759

Training split: 15 of 16

[0] validation\_0-rmse:34076.15009 validation\_1-rmse:30173.91135  
[500] validation\_0-rmse:2186.48824 validation\_1-rmse:2422.19548  
[813] validation\_0-rmse:2021.39231 validation\_1-rmse:2313.93976

Training split: 16 of 16

[0] validation\_0-rmse:33939.52789 validation\_1-rmse:29495.56077  
[500] validation\_0-rmse:2189.84373 validation\_1-rmse:2304.33924  
[977] validation\_0-rmse:1978.42327 validation\_1-rmse:2270.01215

Overall Metrics for Accuracy and Model Performance

Resulting Metrics	
Average RMSE	2,411.87
Average Absolute Accuracy	93.83%
Average Relative Accuracy	94.32%
Average Sharpe Ratio	0.29

Predictions vs Actual Targets

head(5)						tail(5)					
	actual	prediction	abs_acc	rel_acc	sharpe		actual	prediction	abs_acc	rel_acc	sharpe
0	41,297	43,600.54	94.72	92.88	0.36	140155	25,634	28,081.06	91.29	92.44	0.38
1	44,813	45,461.97	98.57	98.00	0.10	140156	24,788	26,979.71	91.88	93.23	0.34
2	46,139	46,121.54	99.96	99.95	0.00	140157	24,365	26,686.66	91.30	92.83	0.36
3	45,716	45,878.23	99.65	99.50	0.03	140158	24,766	25,165.46	98.41	98.77	0.06
4	45,342	45,525.46	99.60	99.43	0.03	140159	24,843	24,667.06	99.29	99.46	0.03

## Predicting into the Future

### 1. Retrain Model with All Training Data

```
df = data.copy()
```

```
def future_predicting_model(df, target_column,
                            lags = False,
                            lag_label_list = None,
                            base_score = 0.5,
                            booster = 'gbtree',
                            n_estimators = 500,
                            early_stopping_rounds = 50,
                            objective = 'reg:squarederror',
                            max_depth = 3,
                            learning_rate = 0.012,
                            verbose = 250):

    from sklearn.metrics import mean_squared_error
    from sklearn.model_selection import TimeSeriesSplit

    # ..... Establishing Features & Targets ..... #

    df = featurize_datetime_index(df)
    target = target_column

    # ..... Lag Features & Feature Cleanup ..... #

    if lags:
        df = year_lags(df, target_column, lag_label_list)
        features = list(df.columns)
        features.remove('weekday_name')
        features.remove('month_name')
        features.remove('time_of_day')
        features.remove(target)
    else:
        features = list(df.columns)
        features.remove('weekday_name')
        features.remove('month_name')
        features.remove('time_of_day')
        features.remove(target)

    # ..... Correcting Datatypes for Model ..... #

    df['week_of_year'] = df.week_of_year.astype(float)

    input_data = df[features]
```

```

target_data = df[target_column]

model = XGBRegressor(base_score = base_score,
                     booster = booster,
                     n_estimators = n_estimators,
                     early_stopping_rounds = early_stopping_rounds,
                     objective = objective,
                     max_depth = max_depth,
                     learning_rate = learning_rate)

trained_model = model.fit(input_data, target_data,
                          eval_set = [(input_data, target_data)],
                          verbose = verbose)

return trained_model

```

```

trained_model = future_predicting_model(df,
                                       'consumption',
                                       lags = True,
                                       lag_label_list = lag_label_list)

```

```

[0] validation_0-rmse:33789.62223
[250] validation_0-rmse:3052.85375
[499] validation_0-rmse:2203.62563

```

## 2. Create an empty dataframe for the date range to predict

- Starting from the end of the input data
- Predicting almost 1 year into the future
- Frequency will be 1 hour, just like the input data frequency

```
df.index.max()
```

```
Timestamp('2022-12-31 23:30:00')
```

```
prediction_range = pd.date_range('2023-01-01', '2024-01-01',
                                freq = '1h')
```

```
see(prediction_range[0:6], "First rows of prediction df")
```

First rows of prediction df

	0
2023-01-01 00:00:00	2023-01-01 00:00:00
2023-01-01 01:00:00	2023-01-01 01:00:00
2023-01-01 02:00:00	2023-01-01 02:00:00
2023-01-01 03:00:00	2023-01-01 03:00:00
2023-01-01 04:00:00	2023-01-01 04:00:00

```
future_predictions_df = pd.DataFrame(index = prediction_range)
future_predictions_df['in_future'] = True
orig = df.copy()
orig['in_future'] = False
orig_plus_future = pd.concat([orig, future_predictions_df])
```

### 3. Combining the Past Data with Future

- Adding the datetime and lag features
- Pulling out just the records that need predictions, future\_predictions
- Removing from features the columns the model cannot or should not work with

```
orig_plus_future = featurize_datetime_index(orig_plus_future)
orig_plus_future = year_lags(orig_plus_future,
                             target_column = "consumption",
                             lag_label_list = lag_label_list)
```

```
features = list(orig_plus_future.columns)
features.remove('weekday_name')
features.remove('month_name')
features.remove('time_of_day')
features.remove('consumption')
features.remove('in_future')
```

```
future_predictions = orig_plus_future.query('in_future').copy()
head_tail_horz(future_predictions, 5, 'DF for Future Predictions')
```

DF for Future Predictions

	consumption	holiday	in_future	hour	weekday	weekday_name	month	head(5)			
								month_name	quarter	year	week_of_
2023-01-01	nan	nan	True	0	6	Sunday	1	January	1	2,023	
2023-01-01	nan	nan	True	1	6	Sunday	1	January	1	2,023	
2023-01-01	nan	nan	True	2	6	Sunday	1	January	1	2,023	
2023-01-01	nan	nan	True	3	6	Sunday	1	January	1	2,023	
2023-01-01	nan	nan	True	4	6	Sunday	1	January	1	2,023	

								tail(5)			
	consumption	holiday	in_future	hour	weekday	weekday_name	month	month_name	quarter	year	week_of_
2023-12-31	nan	nan	True	20	6	Sunday	12	December	4	2,023	
2023-12-31	nan	nan	True	21	6	Sunday	12	December	4	2,023	
2023-12-31	nan	nan	True	22	6	Sunday	12	December	4	2,023	
2023-12-31	nan	nan	True	23	6	Sunday	12	December	4	2,023	
2024-01-01	nan	nan	True	0	0	Monday	1	January	1	2,024	

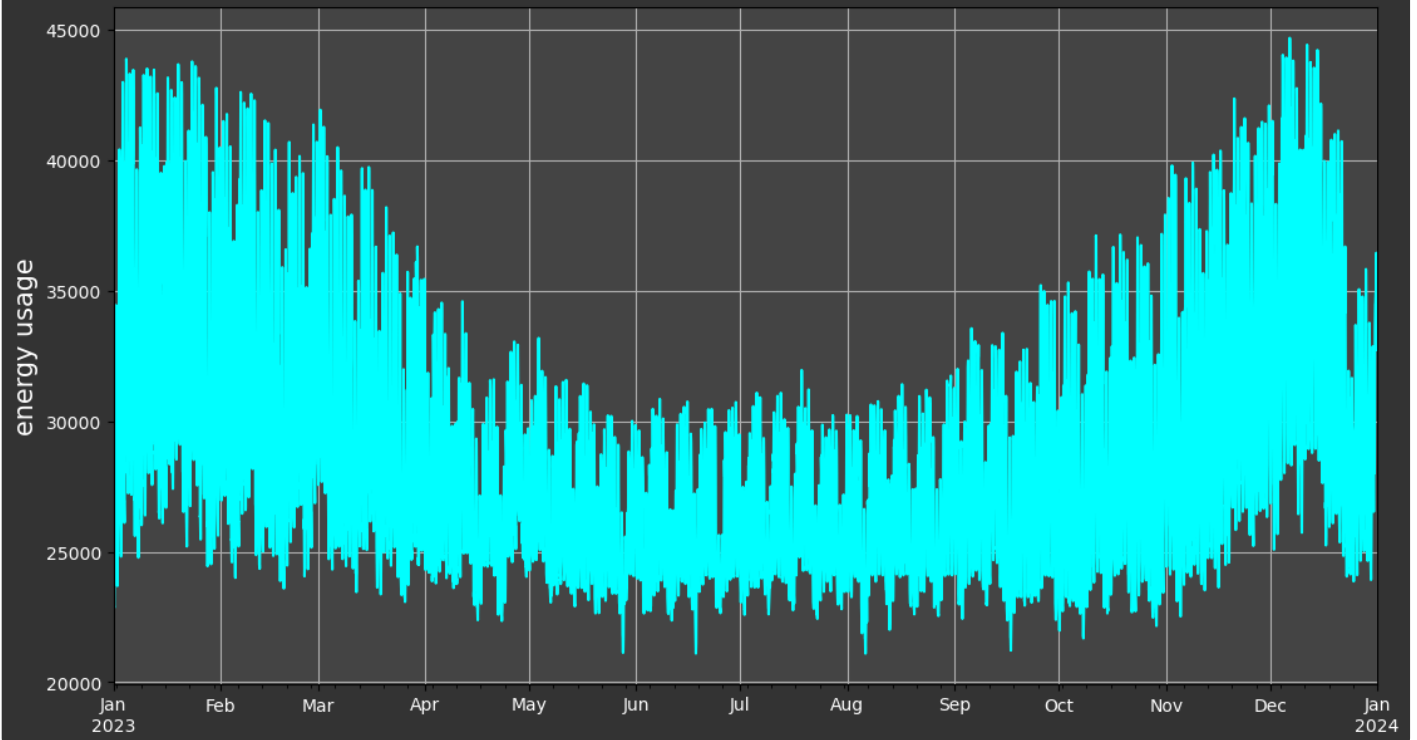
```
future_predictions = future_predictions[features]
future_predictions['week_of_year'] = future_predictions['week_of_year'].astype(float)
```

```
future_predictions['predictions'] = trained_model.predict(future_predictions[features])
```

## 4. Plotting the Prediction Results

```
plt.figure(figsize = (13, 7), facecolor = '#333333')
ax = plt.axes()
future_predictions['predictions'].plot(ax = ax, color = 'cyan')
ax.set_facecolor('#444444')
ax.grid()
plt.xlabel('', color = 'white', fontsize = 15)
plt.ylabel('energy usage', color = 'white', fontsize = 15)
plt.tick_params(labelcolor = 'white', which = 'both')
plt.title('Future Predictions',
          fontsize = 20, pad = 20, color = 'white');
plt.legend().remove()
```

## Future Predictions



---

Sections: ● [Top](#) ● [The Data](#) ● [Feature Engineering](#) ● [Investigating Correlation](#) ● [Lag Features](#) ● [Splitting](#) ● [The Model](#) ● [Results with Traditional Split](#) ● [Using Cross-Validation](#) ● [Making Future Predictions](#)

---